

CSCI 599: Deep Learning and its Applications

Lecture 6

Spring 2019
Shao-Hua Sun

Today's agenda

- **Part 1: Deep Learning Framework**
 - TensorFlow
 - PyTorch
- **Part 2: Cloud Service**
 - Google Cloud
 - Amazon Web Services

Today's agenda

- **Part 1: Deep Learning Framework** (Shao-Hua Sun)
 - TensorFlow (Shao-Hua Sun)
 - PyTorch
- **Part 2: Cloud Service**
 - Google Cloud
 - Amazon Web Services

Today's agenda

- **Part 1: Deep Learning Framework** (Shao-Hua Sun)
 - TensorFlow (Shao-Hua Sun)
 - PyTorch (Te-Lin Wu)
- **Part 2: Cloud Service**
 - Google Cloud
 - Amazon Web Services

Today's agenda

- **Part 1: Deep Learning Framework** (Shao-Hua Sun)
 - TensorFlow (Shao-Hua Sun)
 - PyTorch (Te-Lin Wu)
- **Part 2: Cloud Service** (Hanpeng Liu)
 - Google Cloud
 - Amazon Web Services

Today's agenda

- **Part 1: Deep Learning Framework**
 - TensorFlow
 - PyTorch
- Part 2: Cloud Service
 - Google Cloud
 - Amazon Web Services

Deep Learning Framework

- TensorFlow
- PyTorch
- Torch
- Caffe
- Caffe2
- Theano
- Keras
- Etc.



Deep Learning Framework


Deep Learning Framework



Deep Learning Framework



Deep Learning Framework



Deep Learning Framework



Deep Learning Framework

 torch
NYU/Facebook



PYTORCH
Facebook

theano
U Montreal



Deep Learning Framework

 torch
NYU/Facebook



PYTORCH
Facebook

theano
U Montreal




TensorFlow
Google

Deep Learning Framework

 torch
NYU/Facebook



PYTORCH
Facebook

theano
U Montreal




TensorFlow
Google

Caffe
UC Berkeley

Deep Learning Framework

 torch
NYU/Facebook



PYTORCH
Facebook

theano
U Montreal




TensorFlow
Google

Caffe
UC Berkeley



Deep Learning Framework

 torch
NYU/Facebook



PYTORCH
Facebook

theano
U Montreal




TensorFlow
Google

Caffe
UC Berkeley



 Caffe2
Facebook

Deep Learning Framework

- Model specification
 - Configuration file
 - Caffe
- Programmatic generation
 - PyTorch, Theano, TensorFlow

Deep Learning Framework

- Language
 - C++
 - Caffe
 - Lua
 - Torch
 - Python
 - PyTorch, Theano, TensorFlow

Today's agenda

- **Part 1: Deep Learning Framework**
 - TensorFlow
 - PyTorch
- Part 2: Cloud Service
 - Google Cloud
 - Amazon Web Services

TensorFlow



TensorFlow

- A deep learning library

TensorFlow

- A deep learning library
- Open-sourced by Google

TensorFlow


- A deep learning library
- Open-sourced by Google
- Written with a Python API over a C/C++ engine

TensorFlow

- A deep learning library
- Open-sourced by Google
- Written with a Python API over a C/C++ engine
- Provides primitives for defining functions on tensors and automatically computing their derivatives

TensorFlow

- A deep learning library
- Open-sourced by Google
- Written with a Python API over a C/C++ engine
- Provides primitives for defining functions on tensors and automatically computing their derivatives



Why is it named TensorFlow?

TensorFlow

Tensor

Flow

TensorFlow

- What is a **Tensor**?

TensorFlow

- What is a **Tensor**?
- A **multidimensional array** of numbers

$$T = \begin{array}{|c|c|c|c|c|} \hline & X_{11N} & X_{12N} & X_{13N} & \dots & X_{1NN} \\ \hline X_{111} & X_{112} & X_{113} & \dots & X_{11N} \\ X_{211} & X_{212} & X_{213} & \dots & X_{21N} \\ \vdots & \vdots & \vdots & & \vdots \\ X_{N11} & X_{N12} & X_{N13} & \dots & X_{N1N} \\ \hline \end{array}$$

TensorFlow

- What is a **Tensor**?
- A **multidimensional array** of numbers
- A scalar is a tensor

$$T = \begin{array}{ccccccc} & & & & X_{11N} & X_{12N} & X_{13N} & \dots & X_{1NN} \\ & & & & X_{21N} & X_{22N} & X_{23N} & \dots & X_{2NN} \\ & & & & X_{31N} & X_{32N} & X_{33N} & \dots & X_{3NN} \\ & & & & \vdots & \vdots & \vdots & \ddots & \vdots \\ & & & & X_{N1N} & X_{N2N} & X_{N3N} & \dots & X_{NNN} \end{array}$$

TensorFlow

- What is a **Tensor**?
 - A **multidimensional array** of numbers
- A scalar is a tensor
- A vector is a tensor

$$T = \begin{array}{ccccccc} & & & & X_{11N} & X_{12N} & X_{13N} & \dots & X_{1NN} \\ & & & & X_{21N} & X_{22N} & X_{23N} & \dots & X_{2NN} \\ & & & & X_{31N} & X_{32N} & X_{33N} & \dots & X_{3NN} \\ & & & & \vdots & \vdots & \vdots & \ddots & \vdots \\ & & & & X_{N1N} & X_{N2N} & X_{N3N} & \dots & X_{NNN} \end{array}$$

TensorFlow

- What is a **Tensor**?
 - A **multidimensional array** of numbers
- A scalar is a tensor
- A vector is a tensor
- A matrix is a tensor

$$T = \begin{array}{ccccccc} & & & & X_{11N} & X_{12N} & X_{13N} & \dots & X_{1NN} \\ & & & & X_{21N} & X_{22N} & X_{23N} & \dots & X_{2NN} \\ & & & & X_{31N} & X_{32N} & X_{33N} & \dots & X_{3NN} \\ & & & & \vdots & \vdots & \vdots & \ddots & \vdots \\ & & & & X_{N1N} & X_{N2N} & X_{N3N} & \dots & X_{NNN} \end{array}$$

TensorFlow

- What is a **Tensor**?
 - A **multidimensional array** of numbers
- A scalar is a tensor
- A vector is a tensor
- A matrix is a tensor

$$T = \begin{matrix} & & & & X_{11N} & X_{12N} & X_{13N} & \dots & X_{1NN} \\ & & & & X_{21N} & X_{22N} & X_{23N} & \dots & X_{2NN} \\ & & & & X_{31N} & X_{32N} & X_{33N} & \dots & X_{3NN} \\ & & & & \vdots & \vdots & \vdots & \ddots & \vdots \\ & & & & X_{N1N} & X_{N2N} & X_{N3N} & \dots & X_{NNN} \end{matrix}$$


We use **Tensor** structure to represent **data**

TensorFlow

- Why **Flow**?


TensorFlow

- Why **Flow**?
- Linear classification



TensorFlow

- Why **Flow**?
- Linear classification: can be represented as




TensorFlow

- Why **Flow**?
 - Data flow graph

TensorFlow

- Why **Flow**?
- Data flow graph
- Remember the **computational graph**?




TensorFlow

- Why **Flow**?
 - Data flow graph


TensorFlow

- Why **Flow**?
 - Data flow graph



TensorFlow


- Why **Flow**?
 - Data flow graph



| | |
|-----|------|
| 0.3 | 0.15 |
| 1.5 | 1.91 |


TensorFlow

- Why **Flow**?
 - Data flow graph



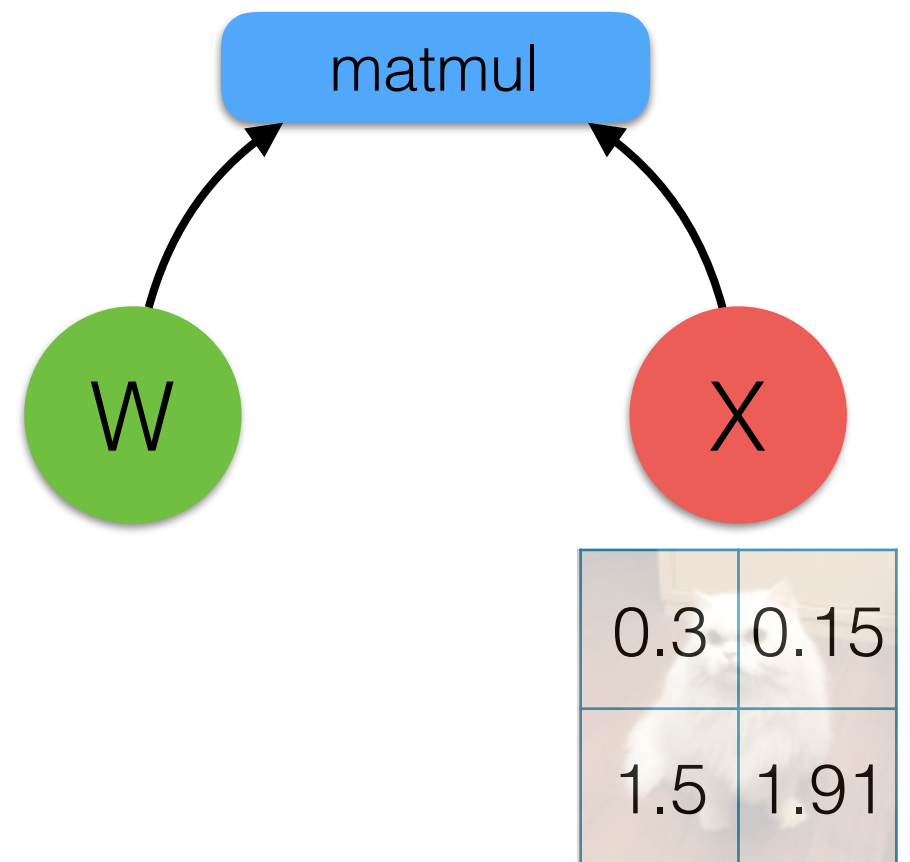
TensorFlow

- Why **Flow**?
 - Data flow graph



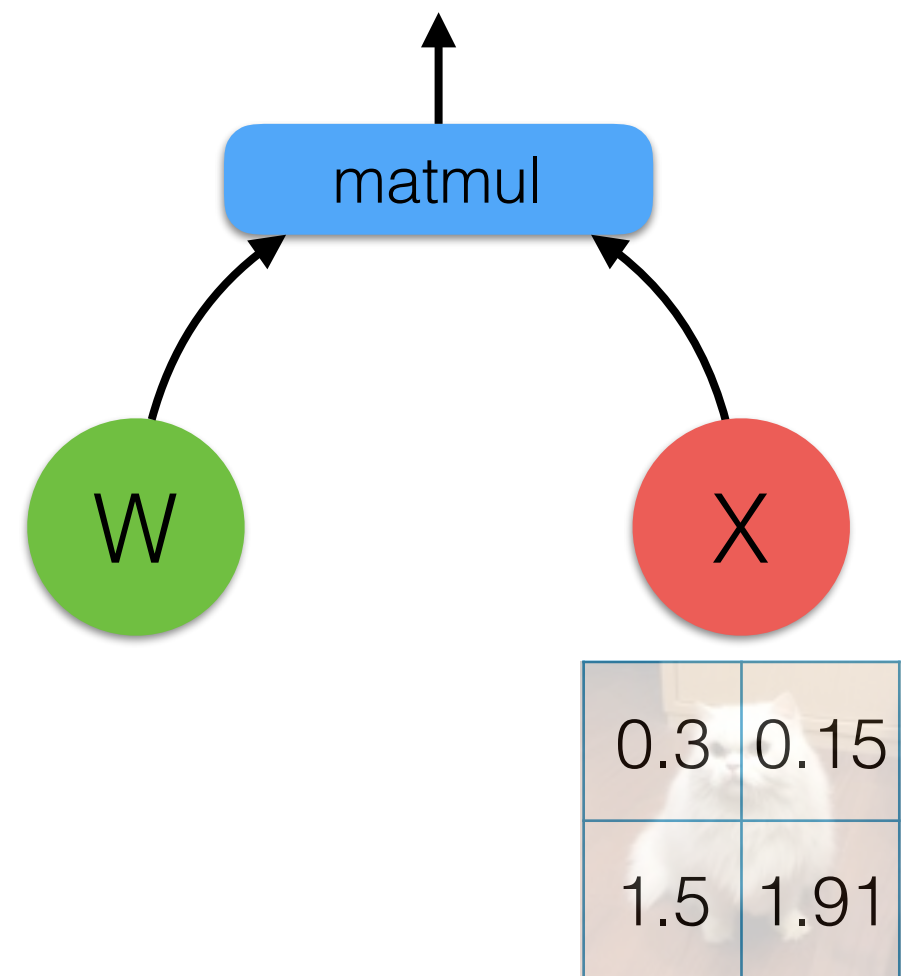
TensorFlow

- Why **Flow**?
 - Data flow graph



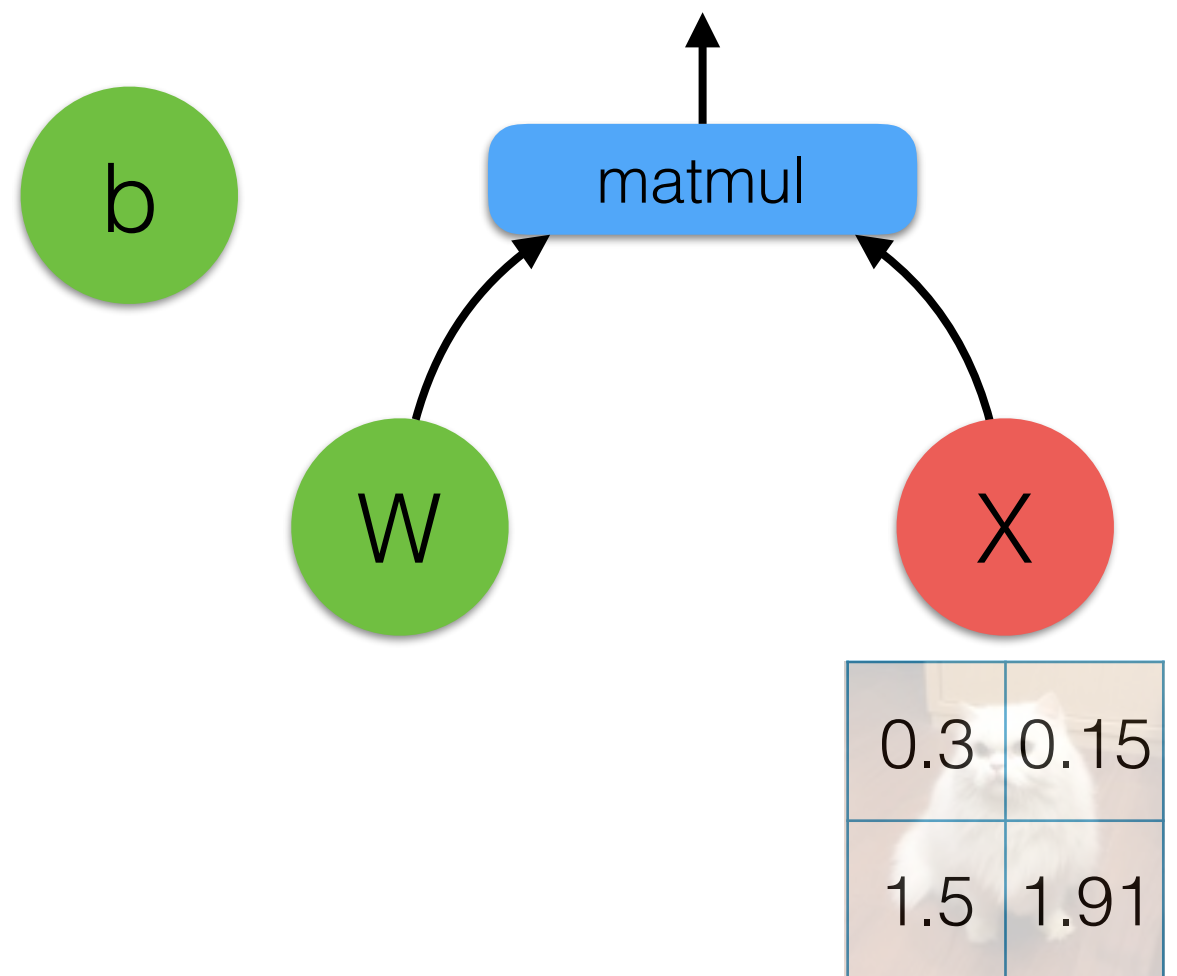
TensorFlow

- Why **Flow**?
 - Data flow graph



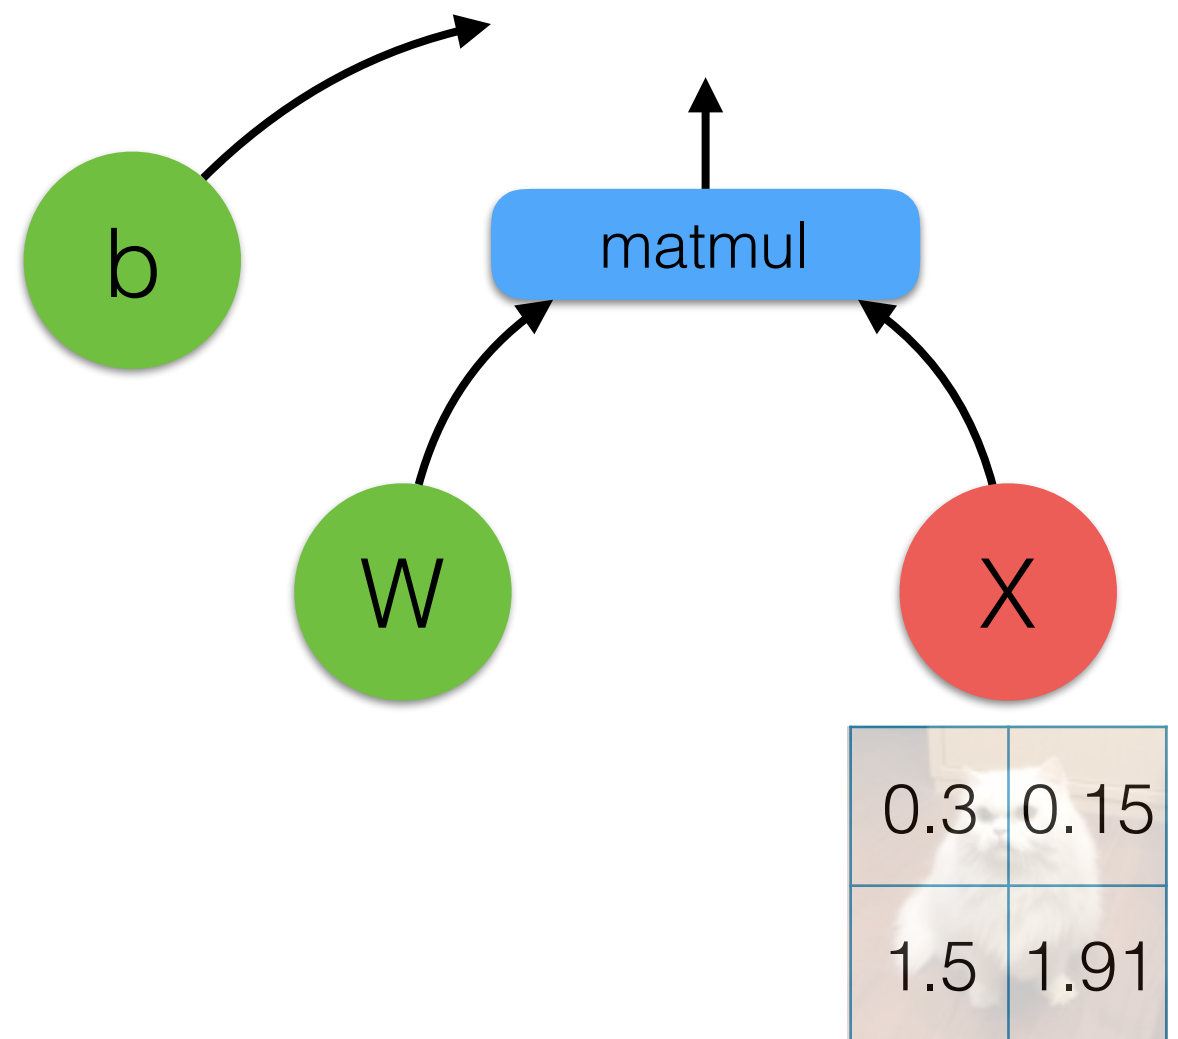
TensorFlow

- Why **Flow**?
 - Data flow graph



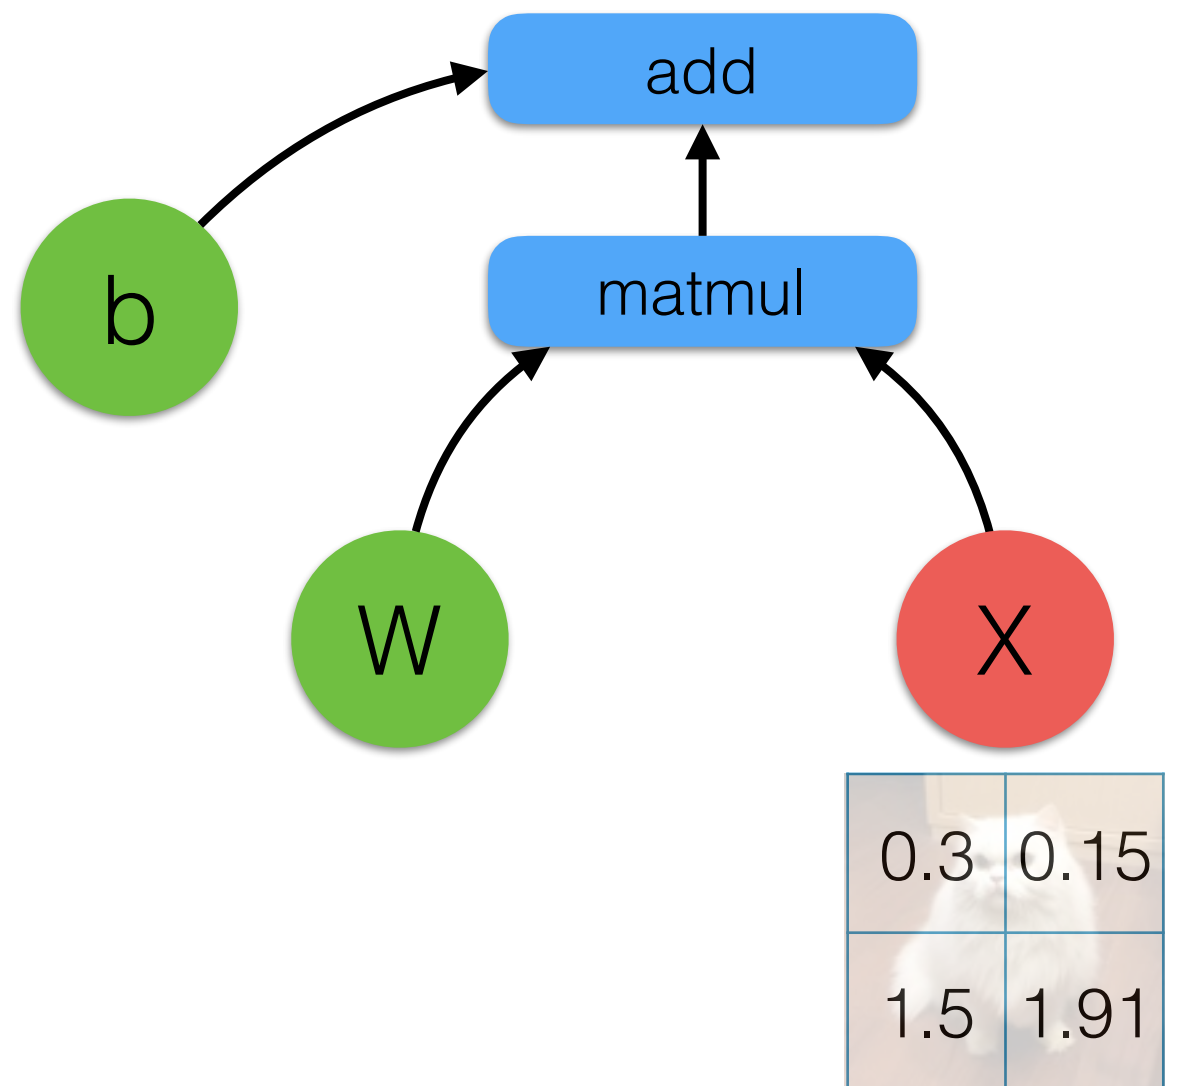
TensorFlow

- Why **Flow**?
 - Data flow graph



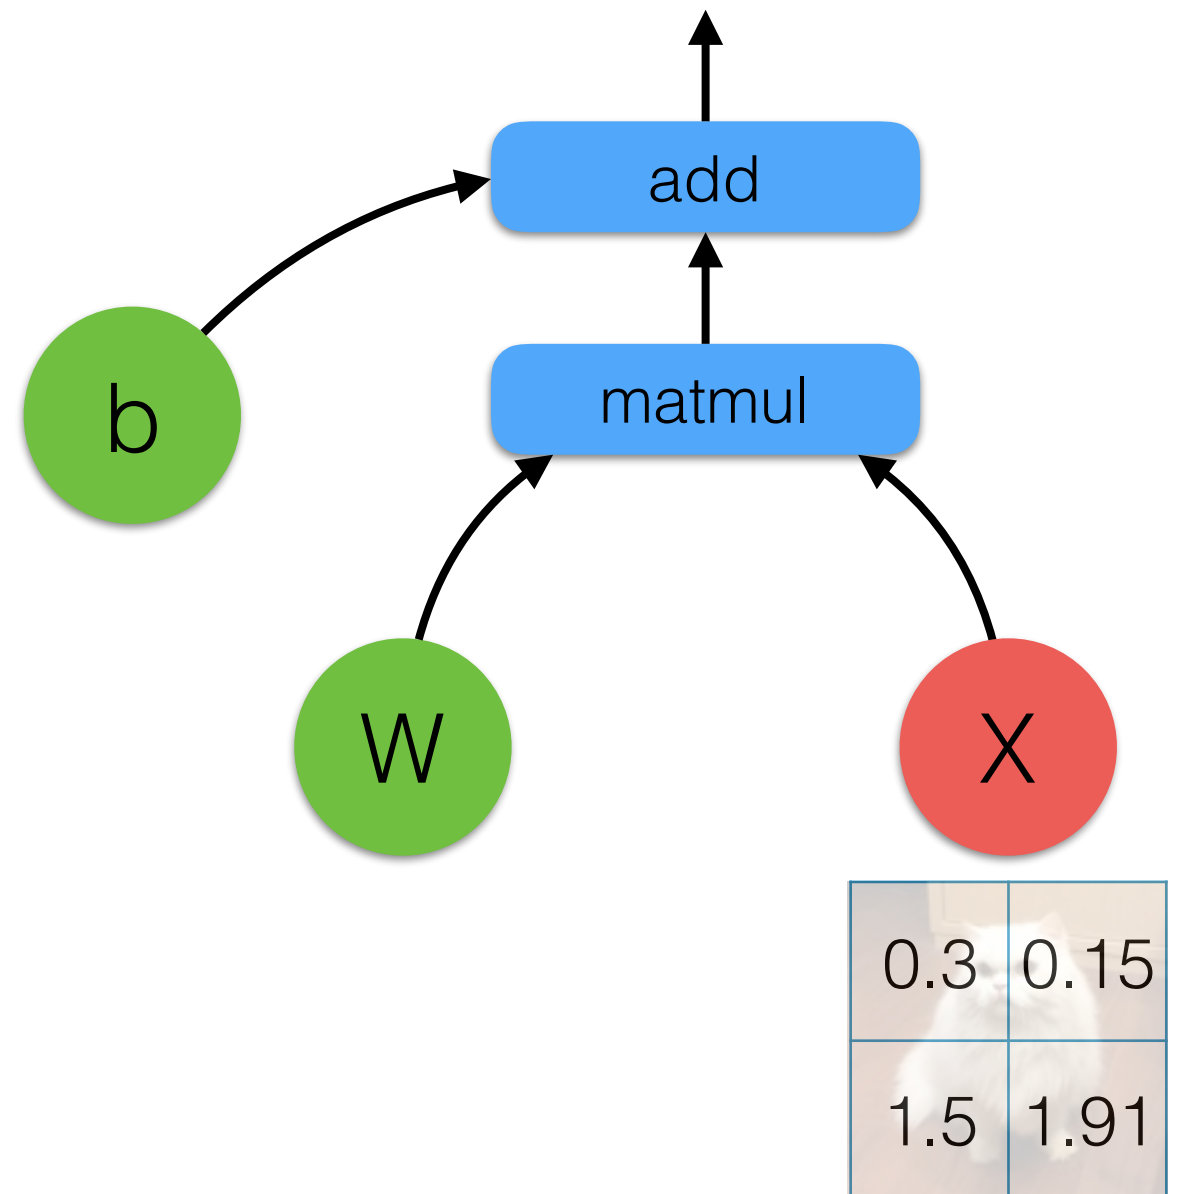
TensorFlow

- Why **Flow**?
 - Data flow graph



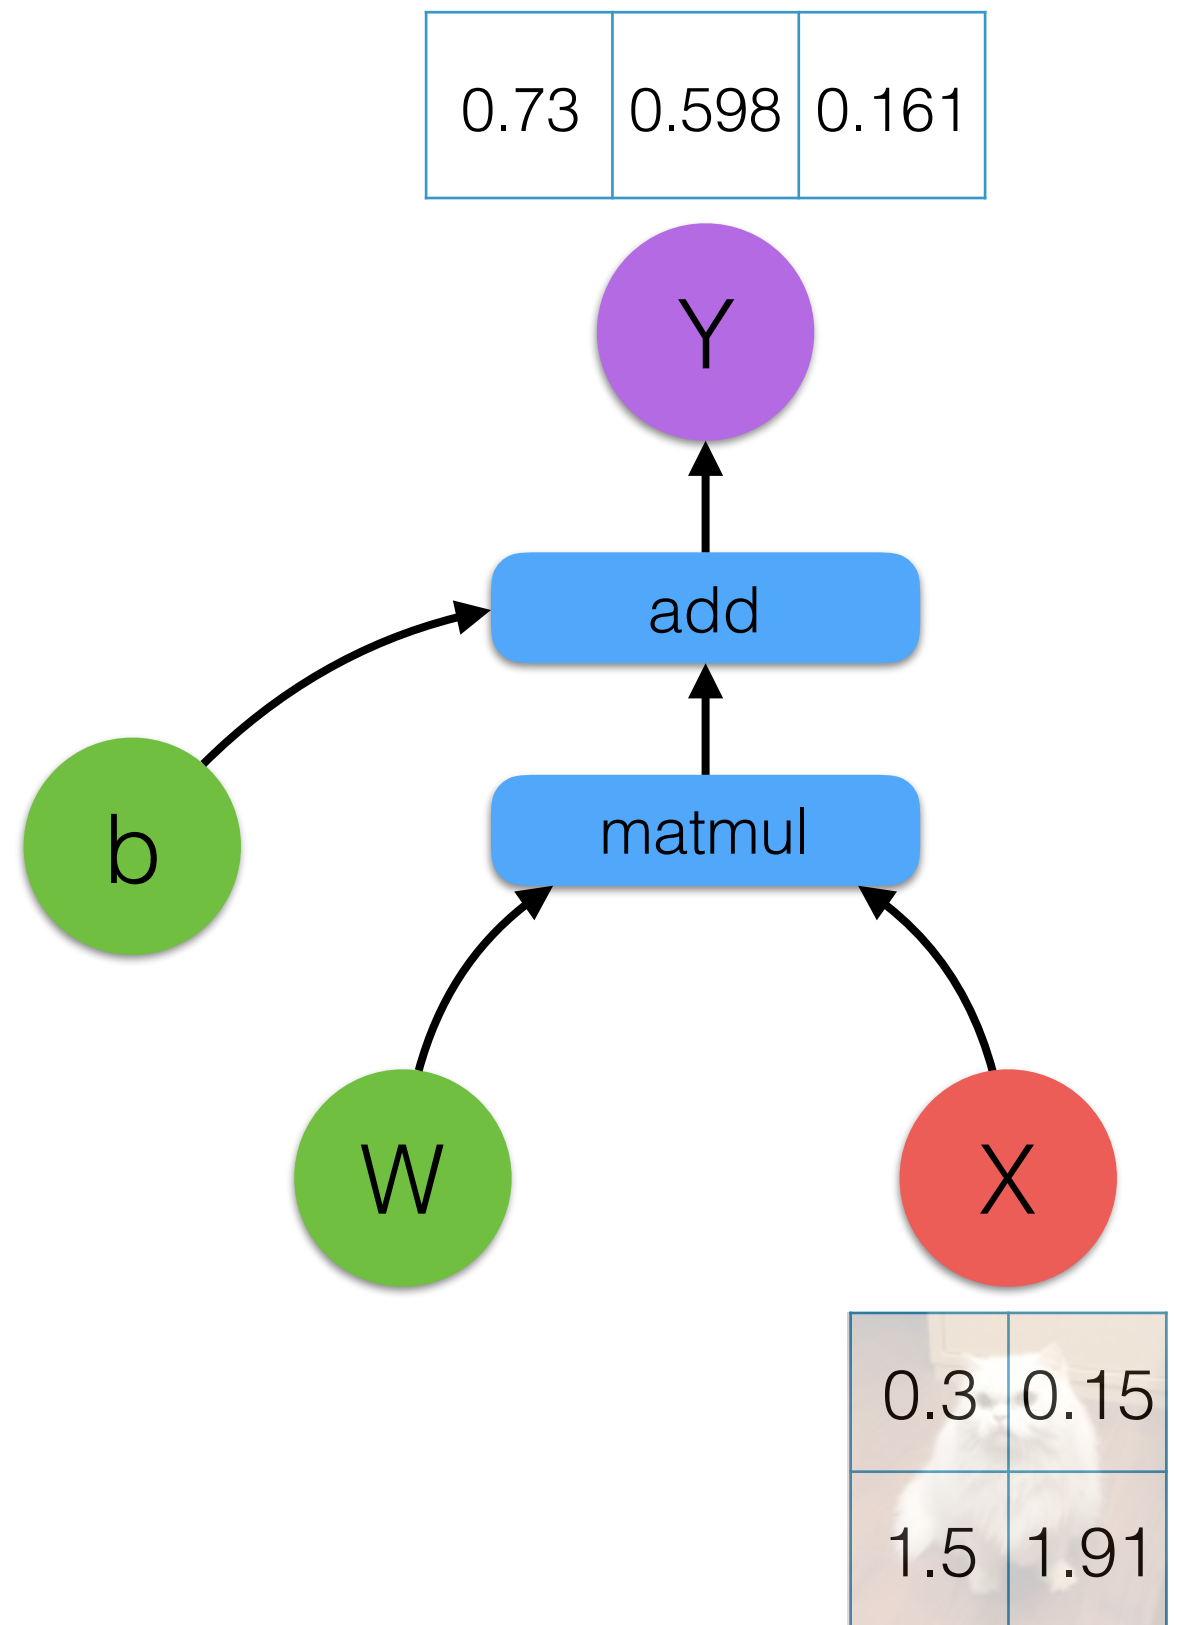
TensorFlow

- Why **Flow**?
 - Data flow graph



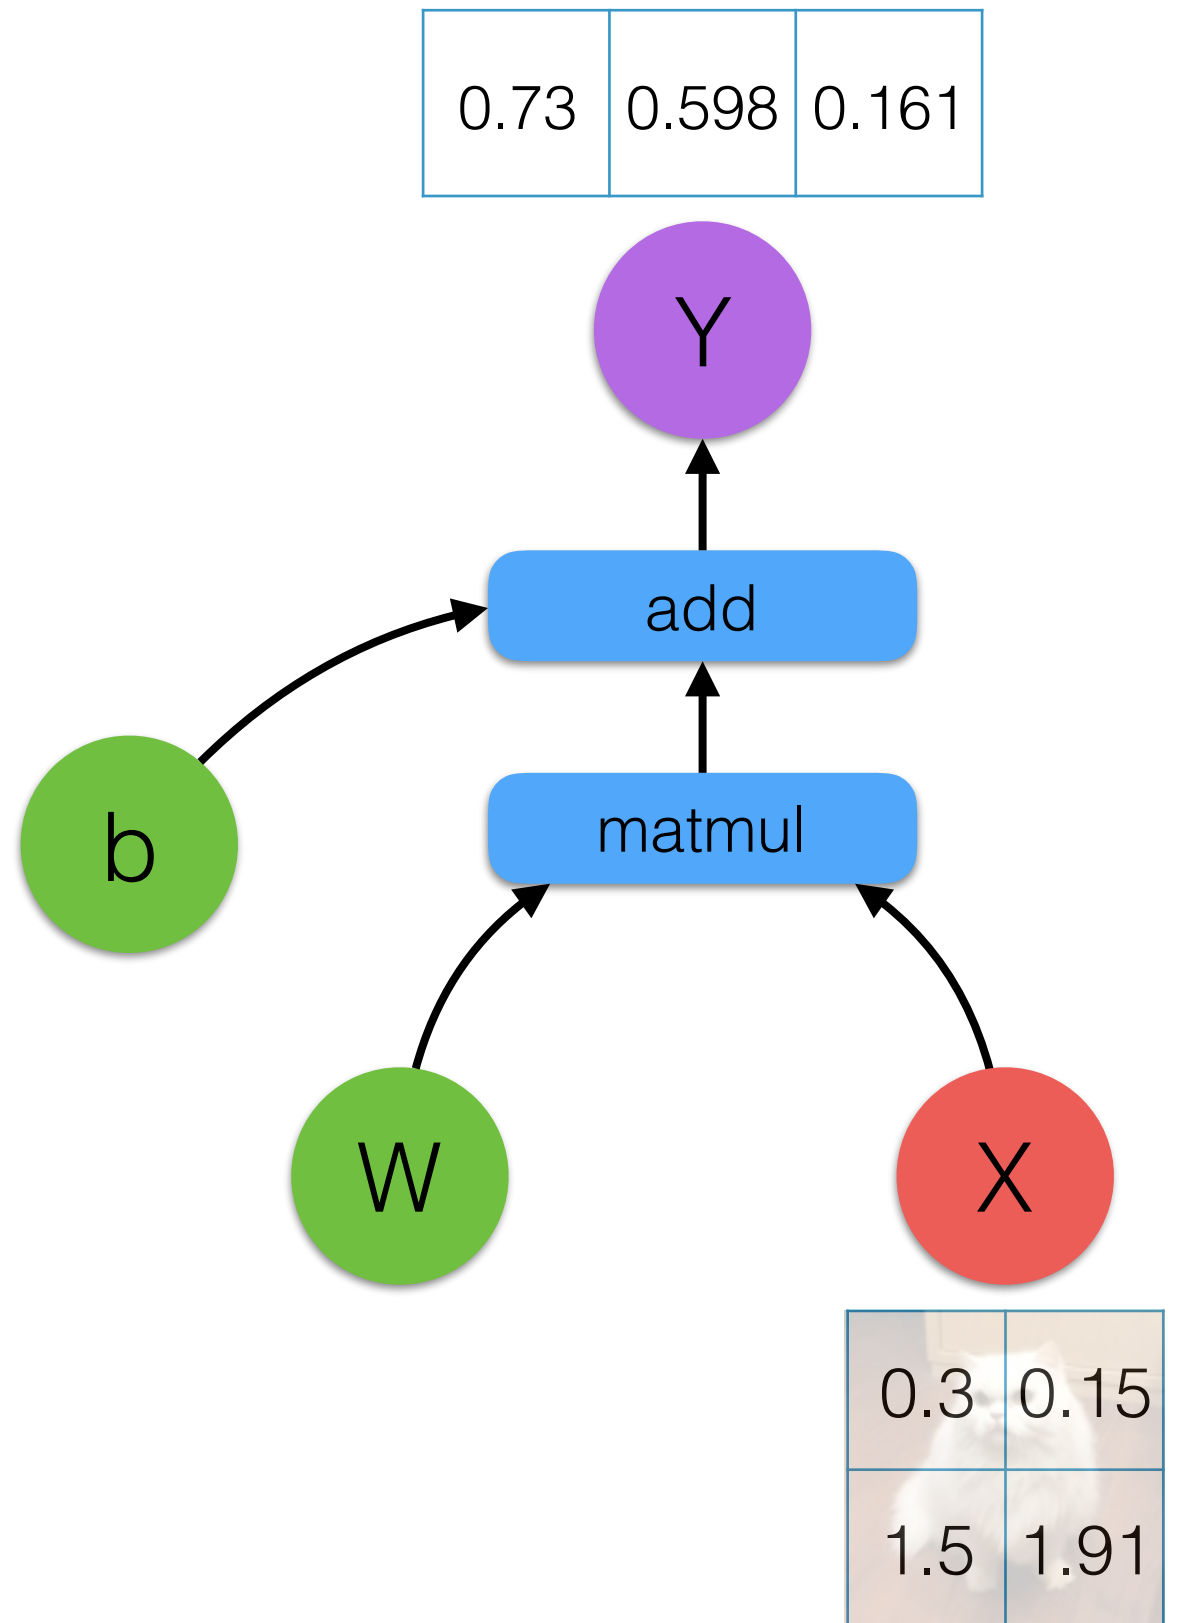
TensorFlow

- Why **Flow**?
- Data flow graph



TensorFlow

- Why **Flow**?
- Data flow graph
- Nodes: operations
- Edges: tensors (data)



Tensors

Flow

Tensors

Flow

Data

Tensors



Data

Flow



Computation

Now you know why it is named TensorFlow

How does it work?

Survey

How much do you know about TensorFlow?

- Never heard about it
- Never used it
- Played with it a little bit
- Worked on a project with it

How does it work?

TensorFlow Core tutorial

This gives Python access to all of TensorFlow's classes, methods, and symbols

```
import tensorflow as tf
```

The Computational Graph

Two sections

- Building the computational graph
- Running the computational graph

The Computational Graph

- Building the computational graph

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Output

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Output

```
Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=float32)
```

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Output

```
Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=float32)
```

- Running the computational graph

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Output

```
Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=float32)
```

- Running the computational graph

```
sess = tf.Session()
print(sess.run([node1, node2]))
```

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Output

```
Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=float32)
```

- Running the computational graph

```
sess = tf.Session()
print(sess.run([node1, node2]))
```

Output

The Computational Graph

- Building the computational graph

```
node1 = tf.constant(3.0, dtype=tf.float32)
node2 = tf.constant(4.0) # also tf.float32 implicitly
print(node1, node2)
```

Output

```
Tensor("Const:0", shape=(), dtype=float32) Tensor("Const_1:0", shape=(), dtype=float32)
```

- Running the computational graph

```
sess = tf.Session()
print(sess.run([node1, node2]))
```

Output

```
[3.0, 4.0]
```

TensorFlow tutorial

Implementation Demo

Implementation Demo

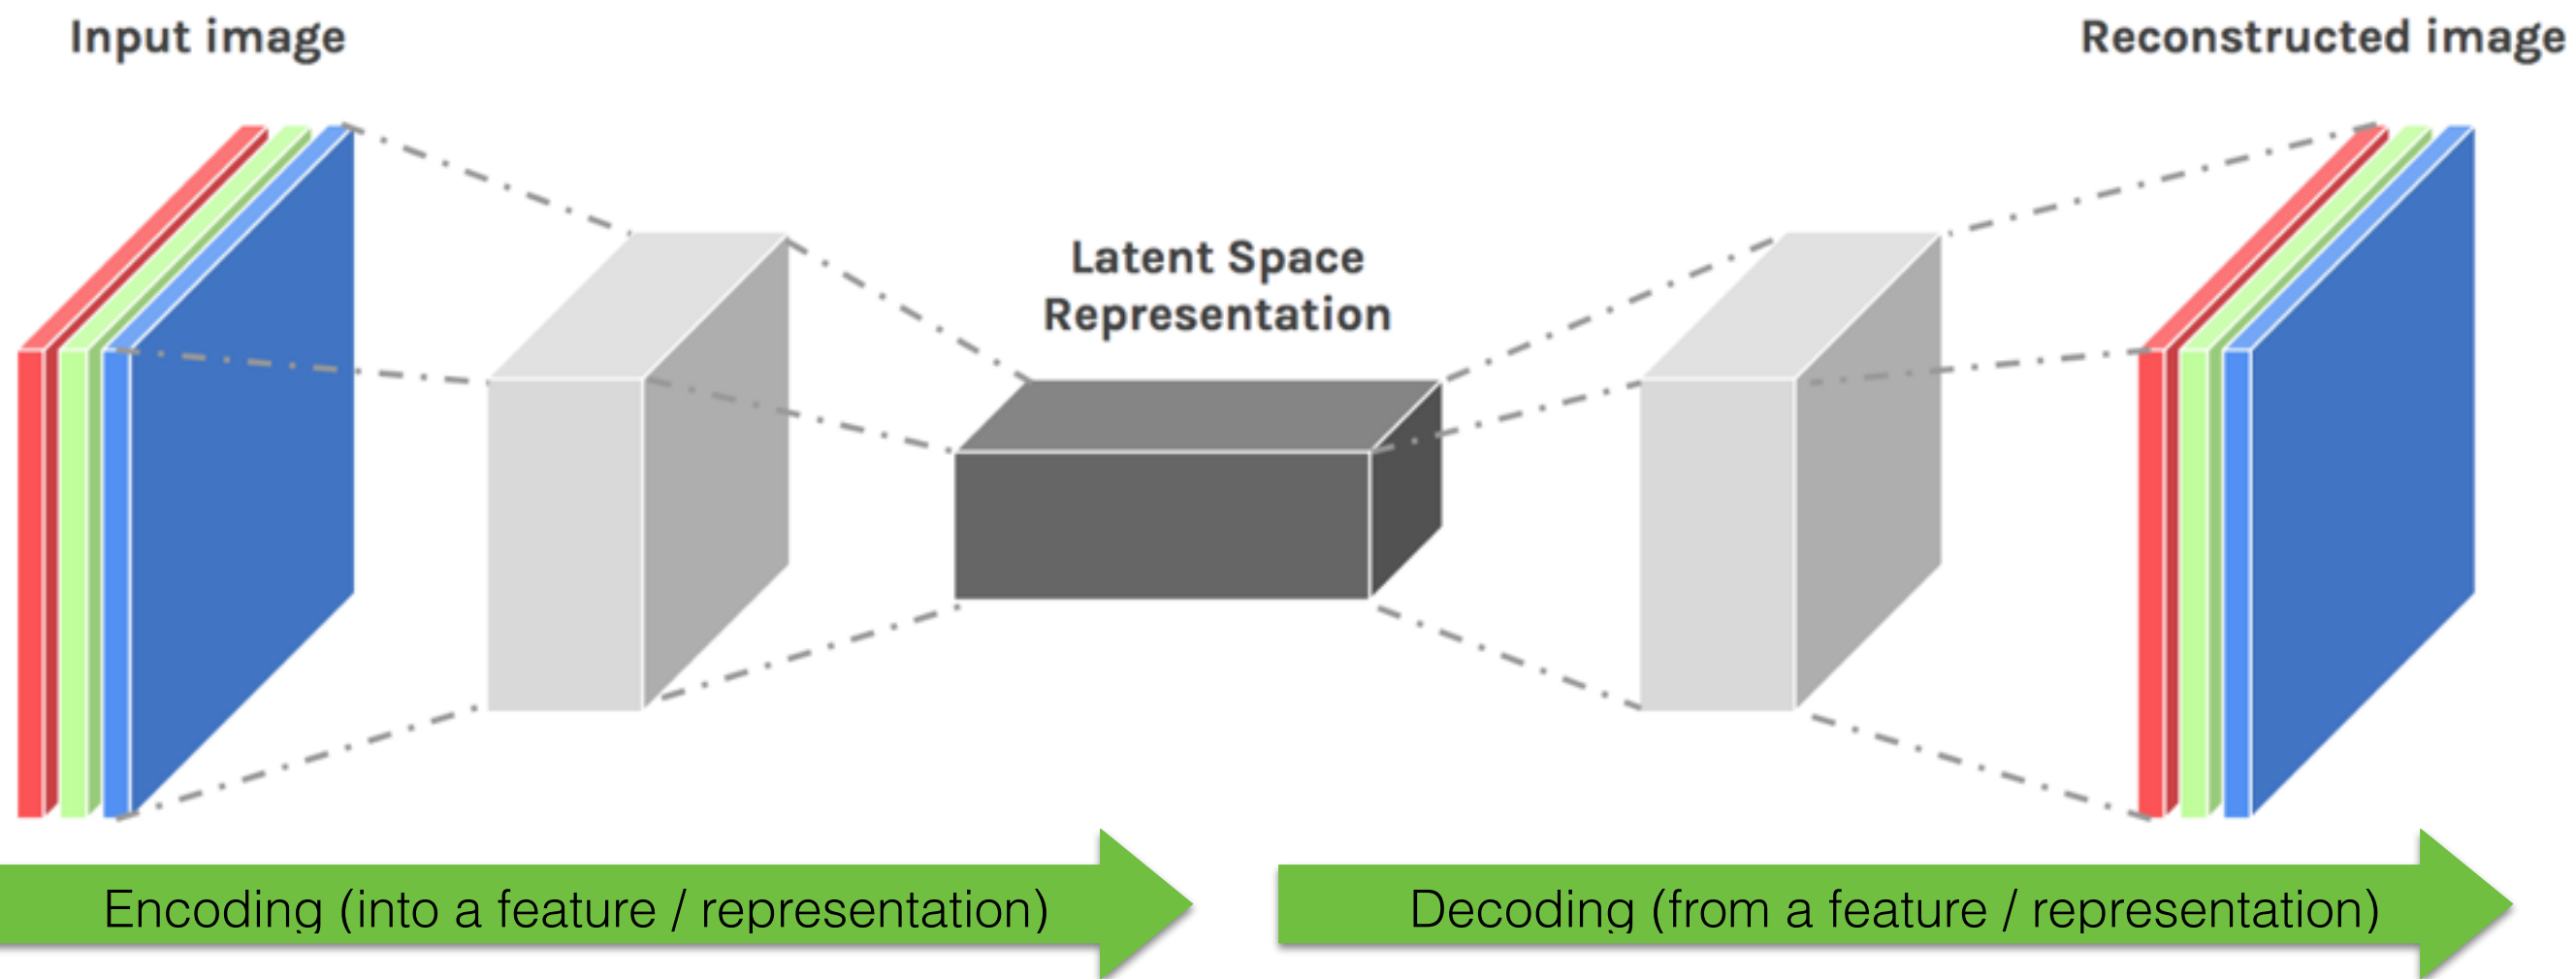
- Building phase
 - Build a network
- Running phase
 - Train the network
 - Test the network

Implementation Demo

- Variational autoencoder
 - Background
 - Model architecture
- Demo
 - Build a network
 - Train the network
 - Test the network

Autoencoder

- Autoencoder




PC: <https://hackernoon.com/latent-space-visualization-deep-learning-bits-2-bd09a46920df>

Autoencoder

- Autoencoder
 - Compression



Autoencoder

- Autoencoder
 - Compression




Autoencoder

- Autoencoder
 - Compression




Autoencoder

- Autoencoder
 - Compression




Autoencoder

- Autoencoder
 - Compression




Autoencoder

- Autoencoder
 - Compression




Autoencoder

- Loss
 - Reconstruction loss




Autoencoder

- Loss
 - Reconstruction loss




Autoencoder

- Model
 - Decoder (3 layers)
 - Encoder (3 layers)




Autoencoder

- Function: training phase
- Train the model given a batch by minimizing the reconstruction loss




Autoencoder

- Function: training phase
- Train the model given a batch by minimizing the reconstruction loss




Autoencoder

- Function: training phase
- Train the model given a batch by minimizing the reconstruction loss




Autoencoder

- Function: testing phase
- Reconstruction



Autoencoder

- Function: testing phase
- Transformation




Coding Session

Autoencoder

Autoencoder

- Results
- Looks pretty good! Doesn't it?



Autoencoder

- What about generation?

Autoencoder

- What about generation?

Images in the dataset



Autoencoder

- What about generation?

Images in the dataset



Autoencoder

- What about generation?

Images in the dataset




Generated images




Autoencoder

- What about generation?




Autoencoder

- What about generation?
- Given feature vectors
- Can we generate data?



Autoencoder

- What about generation?
- Given feature vectors
- Can we generate data?
- No!
- What are the valid feature vectors?




Compressed
Feature Vector





Autoencoder





Autoencoder






Autoencoder




Autoencoder




Autoencoder







Autoencoder







Autoencoder



Autoencoder







Autoencoder



What if...





Autoencoder




What if...








Autoencoder




What if...








Autoencoder




What if...








Autoencoder





What if...








Autoencoder





What if...



Autoencoder




What if...



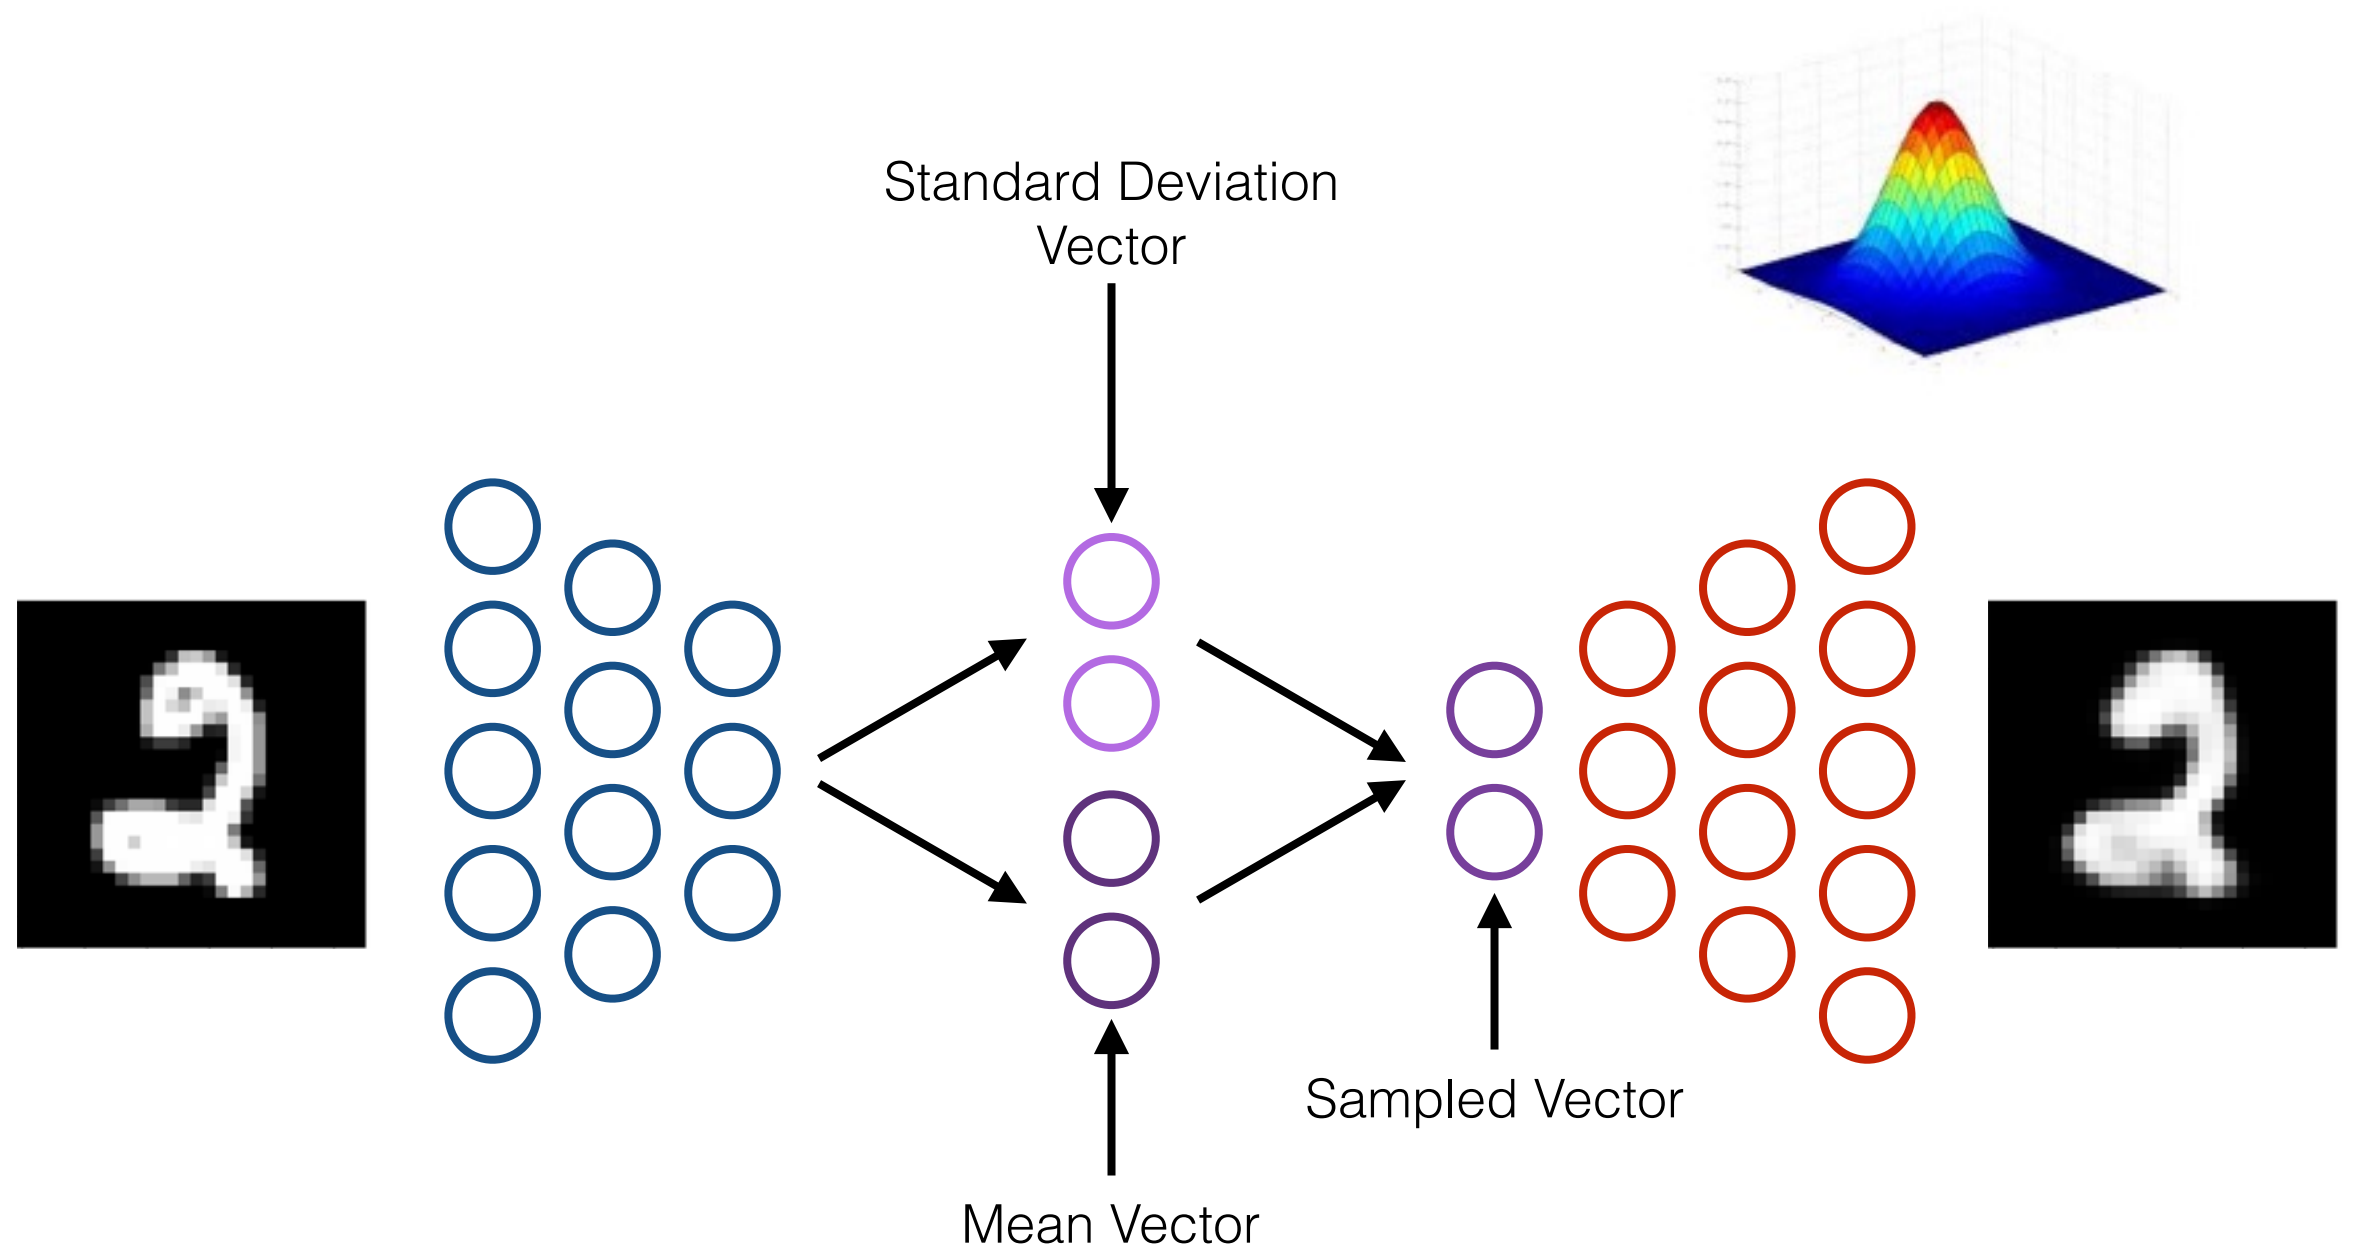
Variational Autoencoder

- Constrain the latent distribution



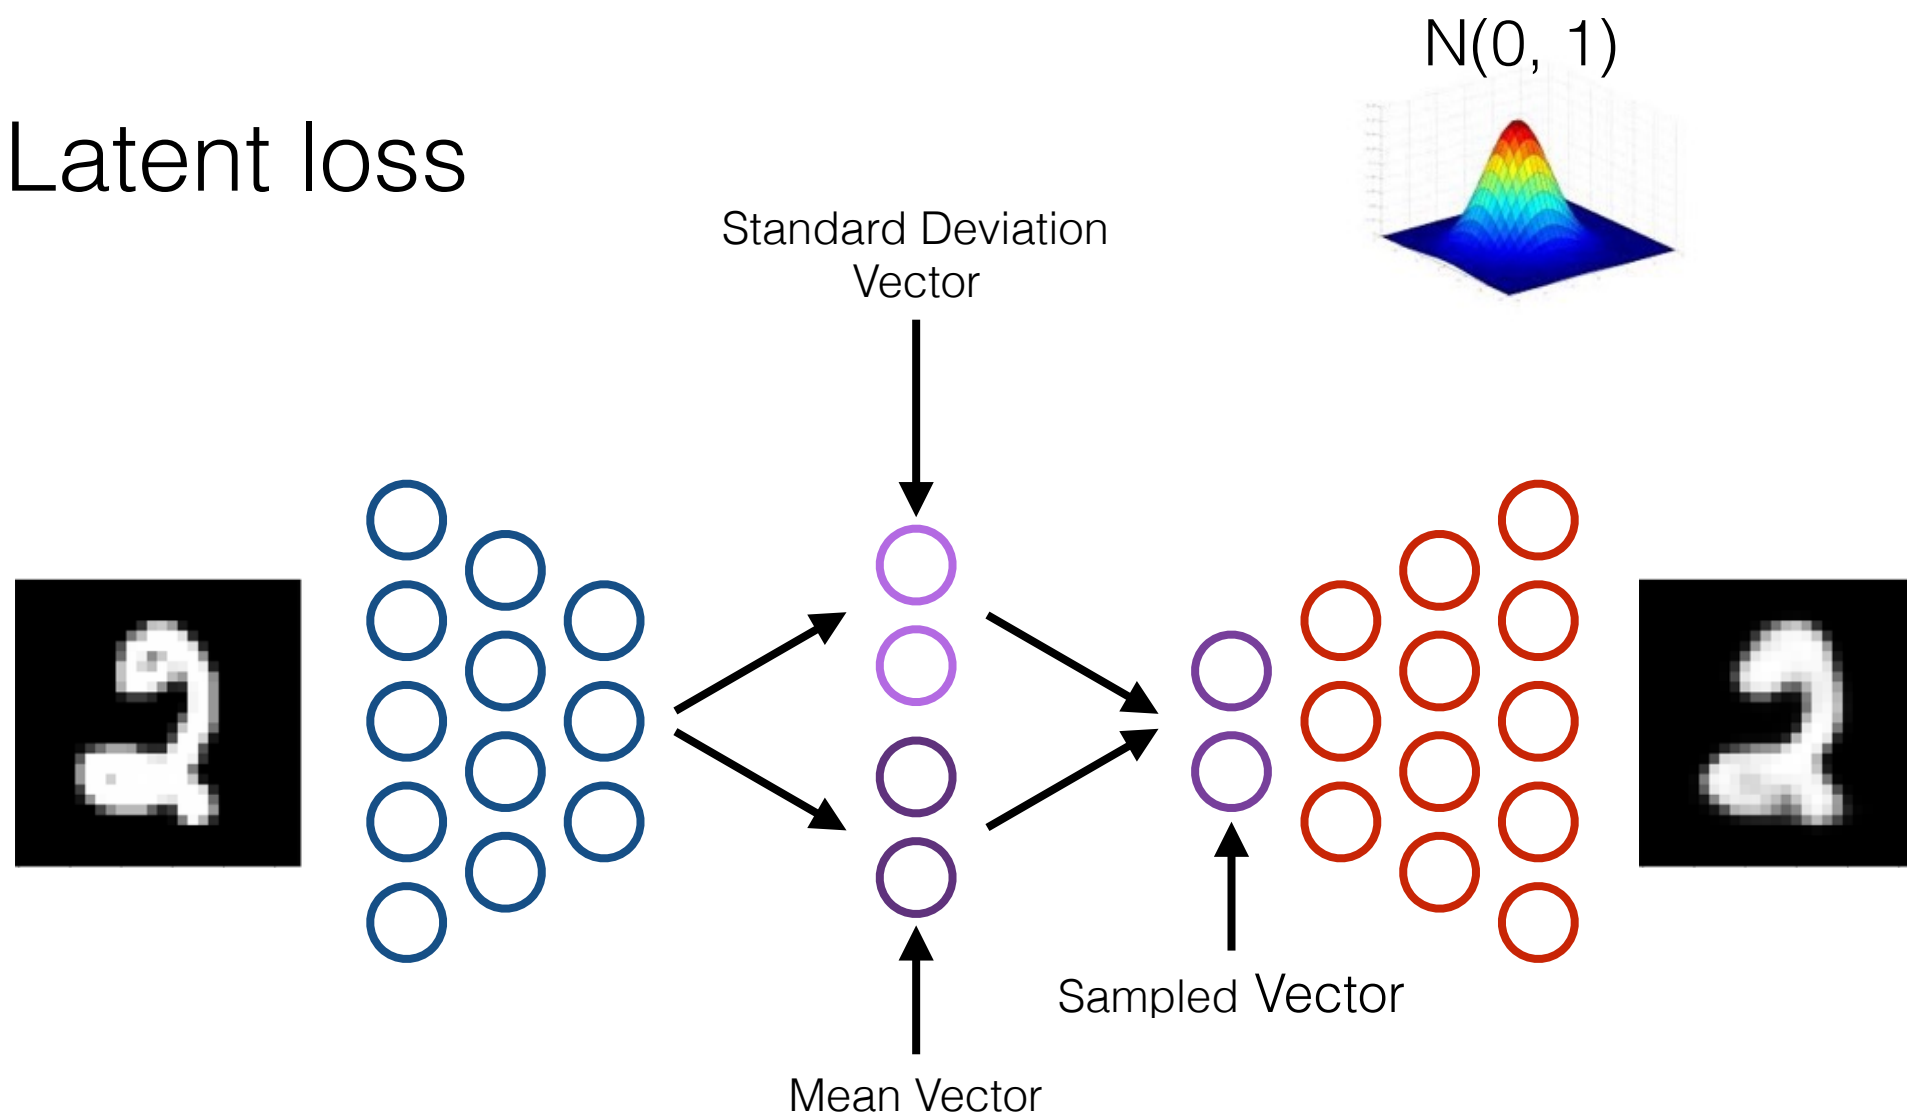
Variational Autoencoder

- Constrain the latent distribution: a Gaussian distribution



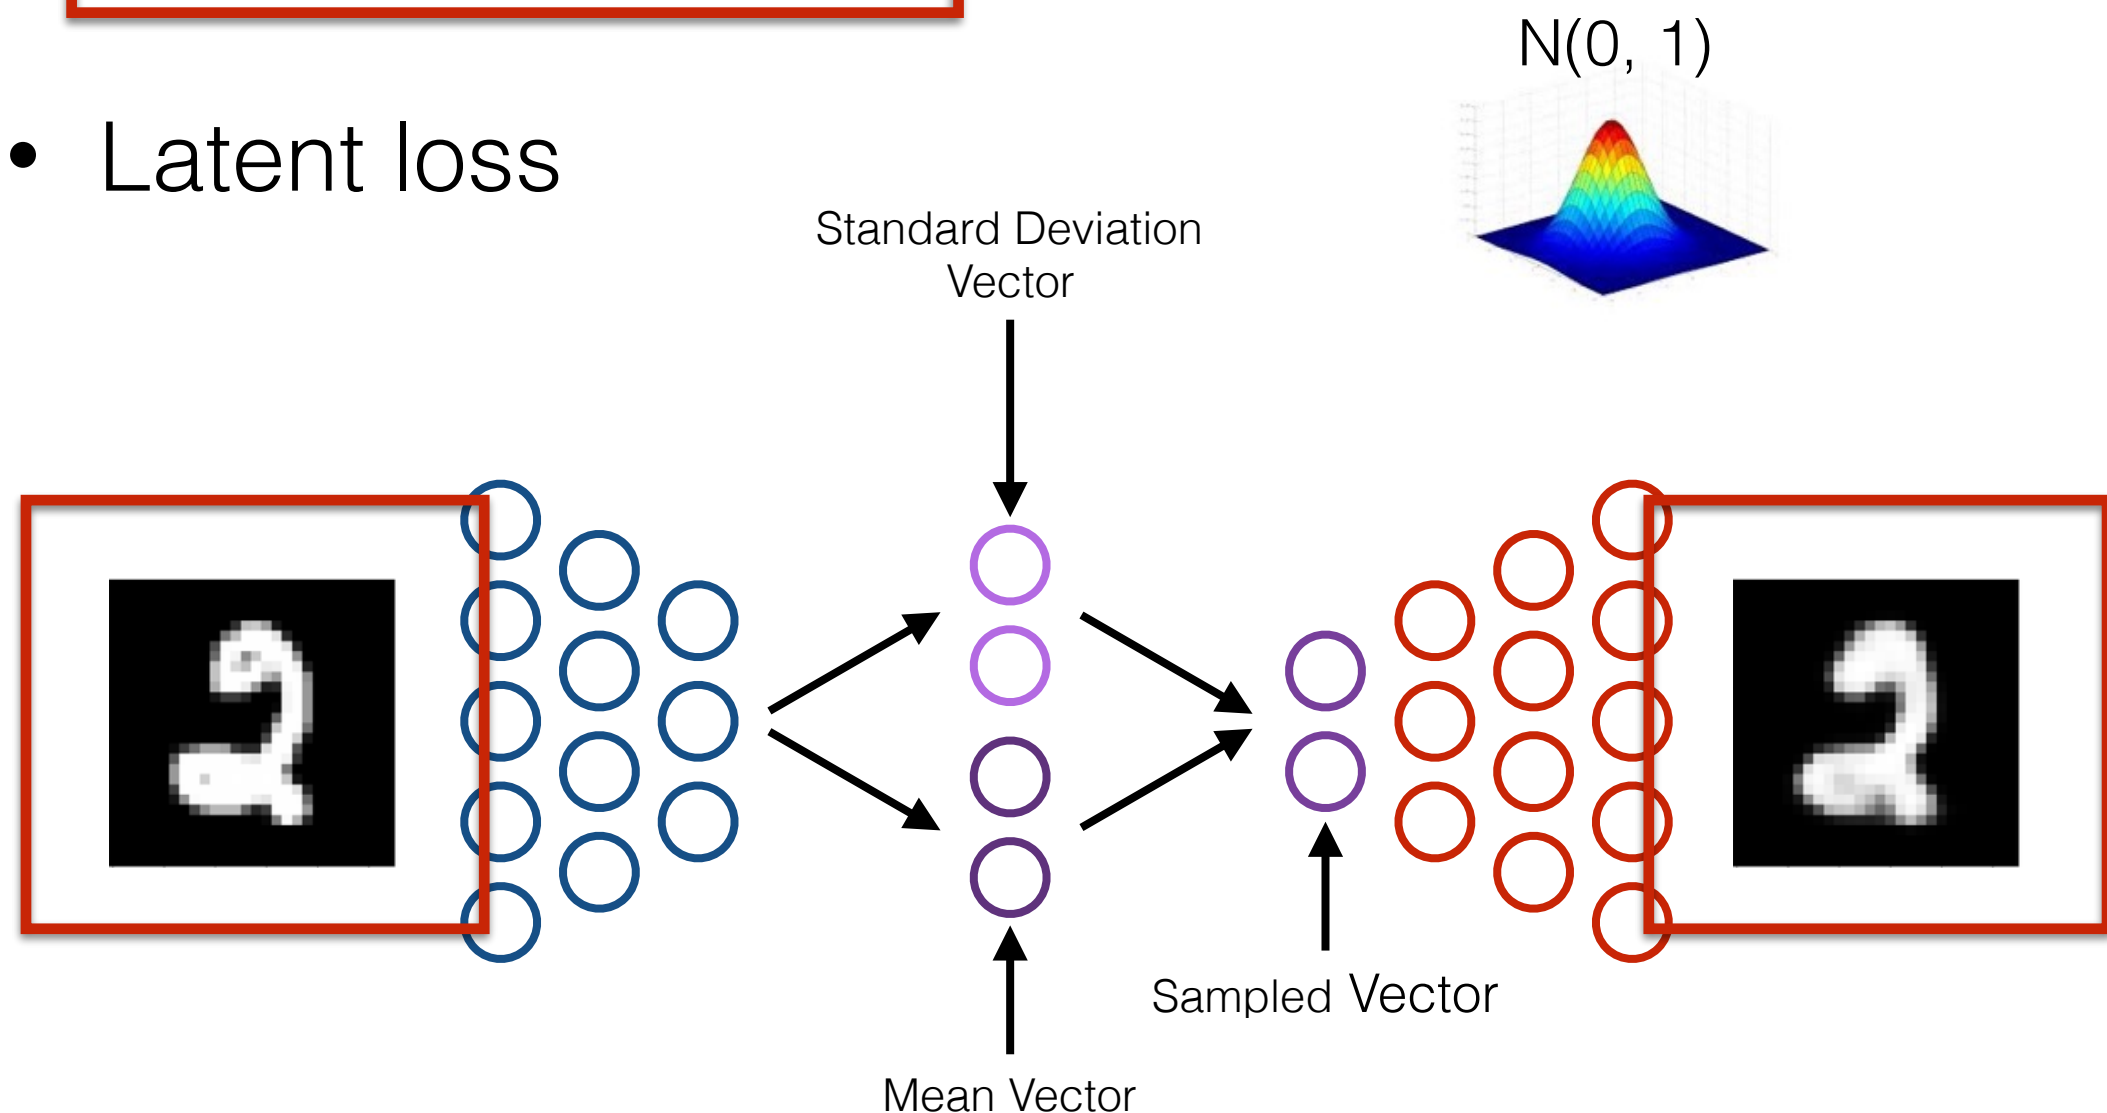
Variational Autoencoder

- Loss
 - Reconstruction loss
 - Latent loss



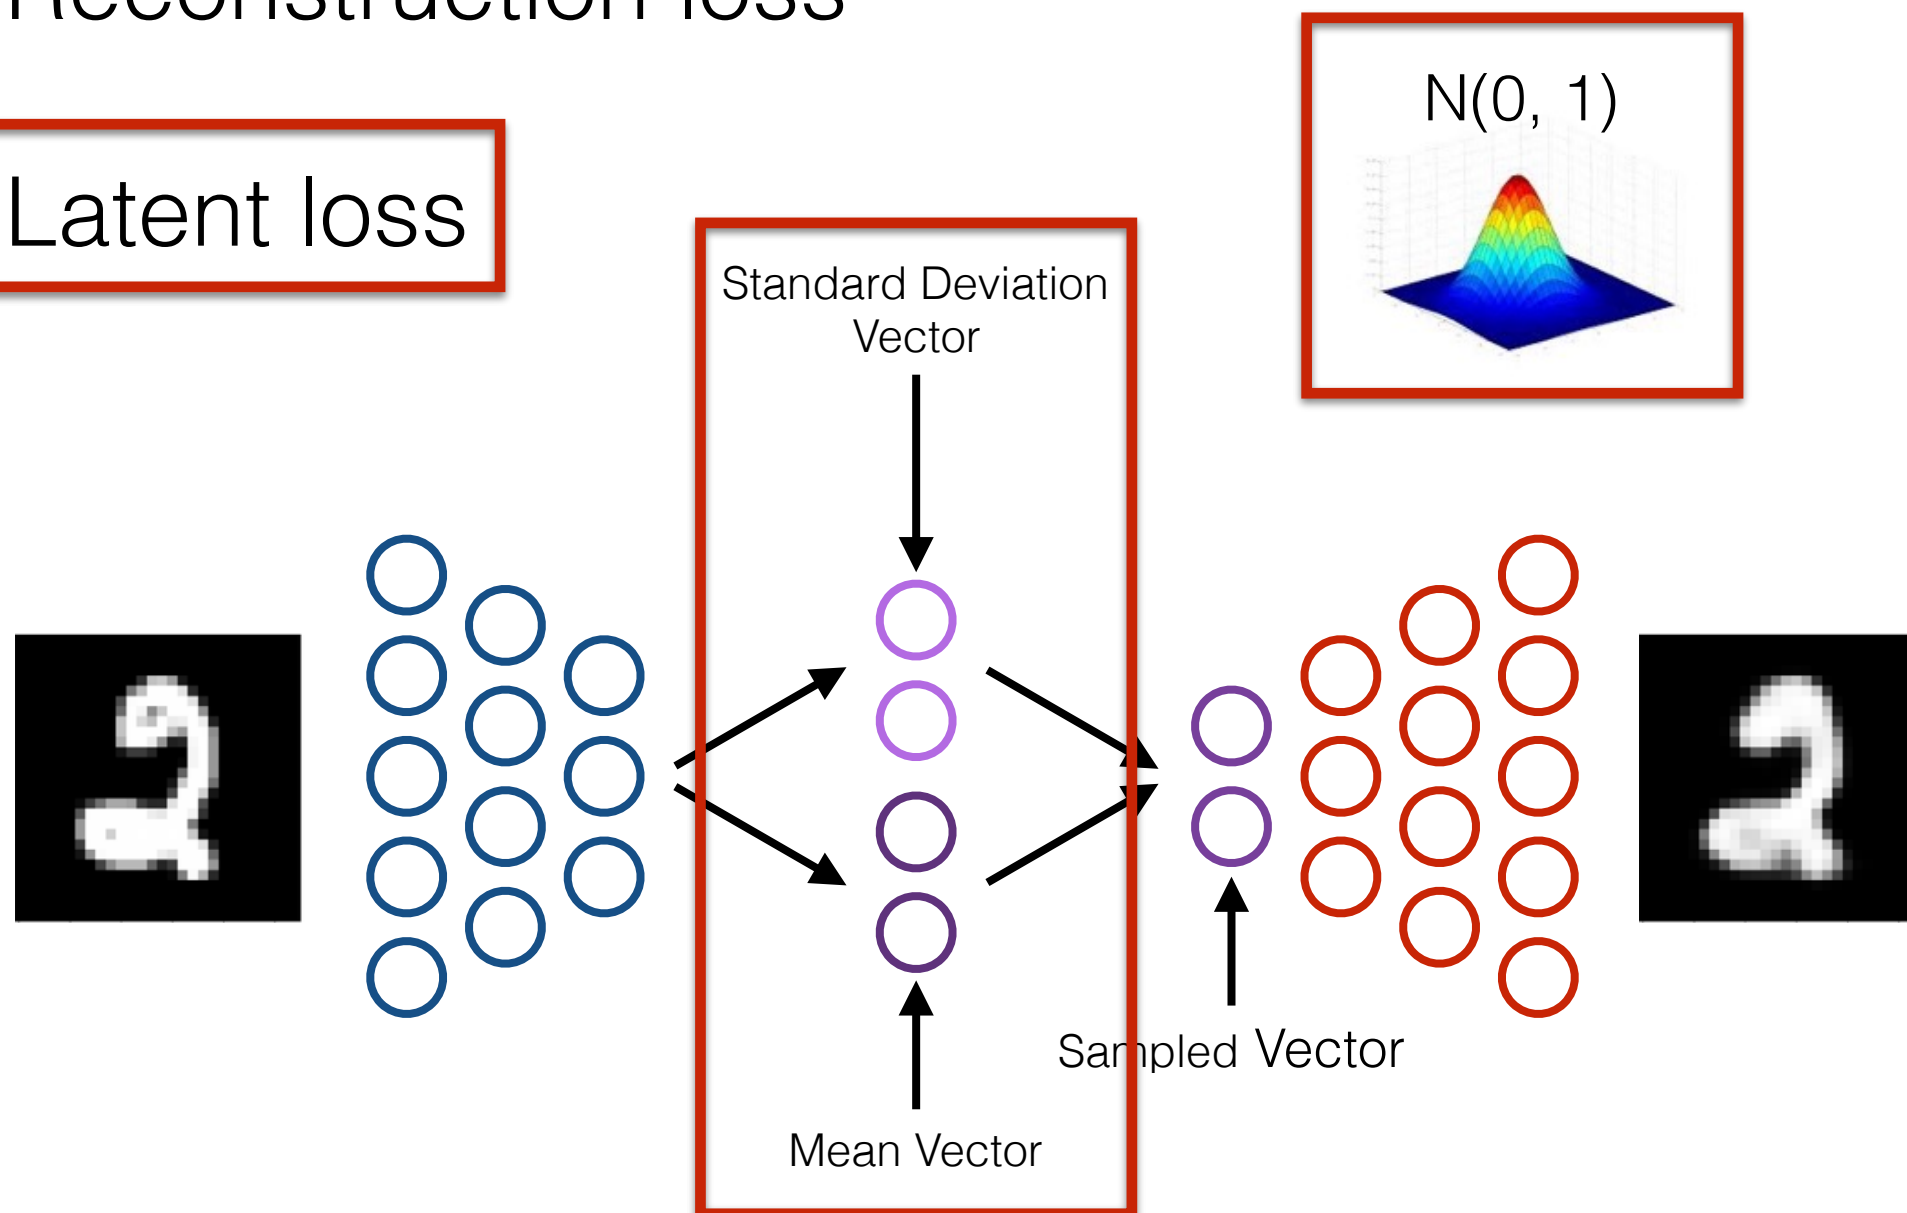
Variational Autoencoder

- Loss
 - Reconstruction loss
- Latent loss




Variational Autoencoder

- Loss
 - Reconstruction loss
 - Latent loss




Variational Autoencoder

- Generation




Variational Autoencoder

- Model
 - Decoder (3 layers)
 - Encoder (3 layers)



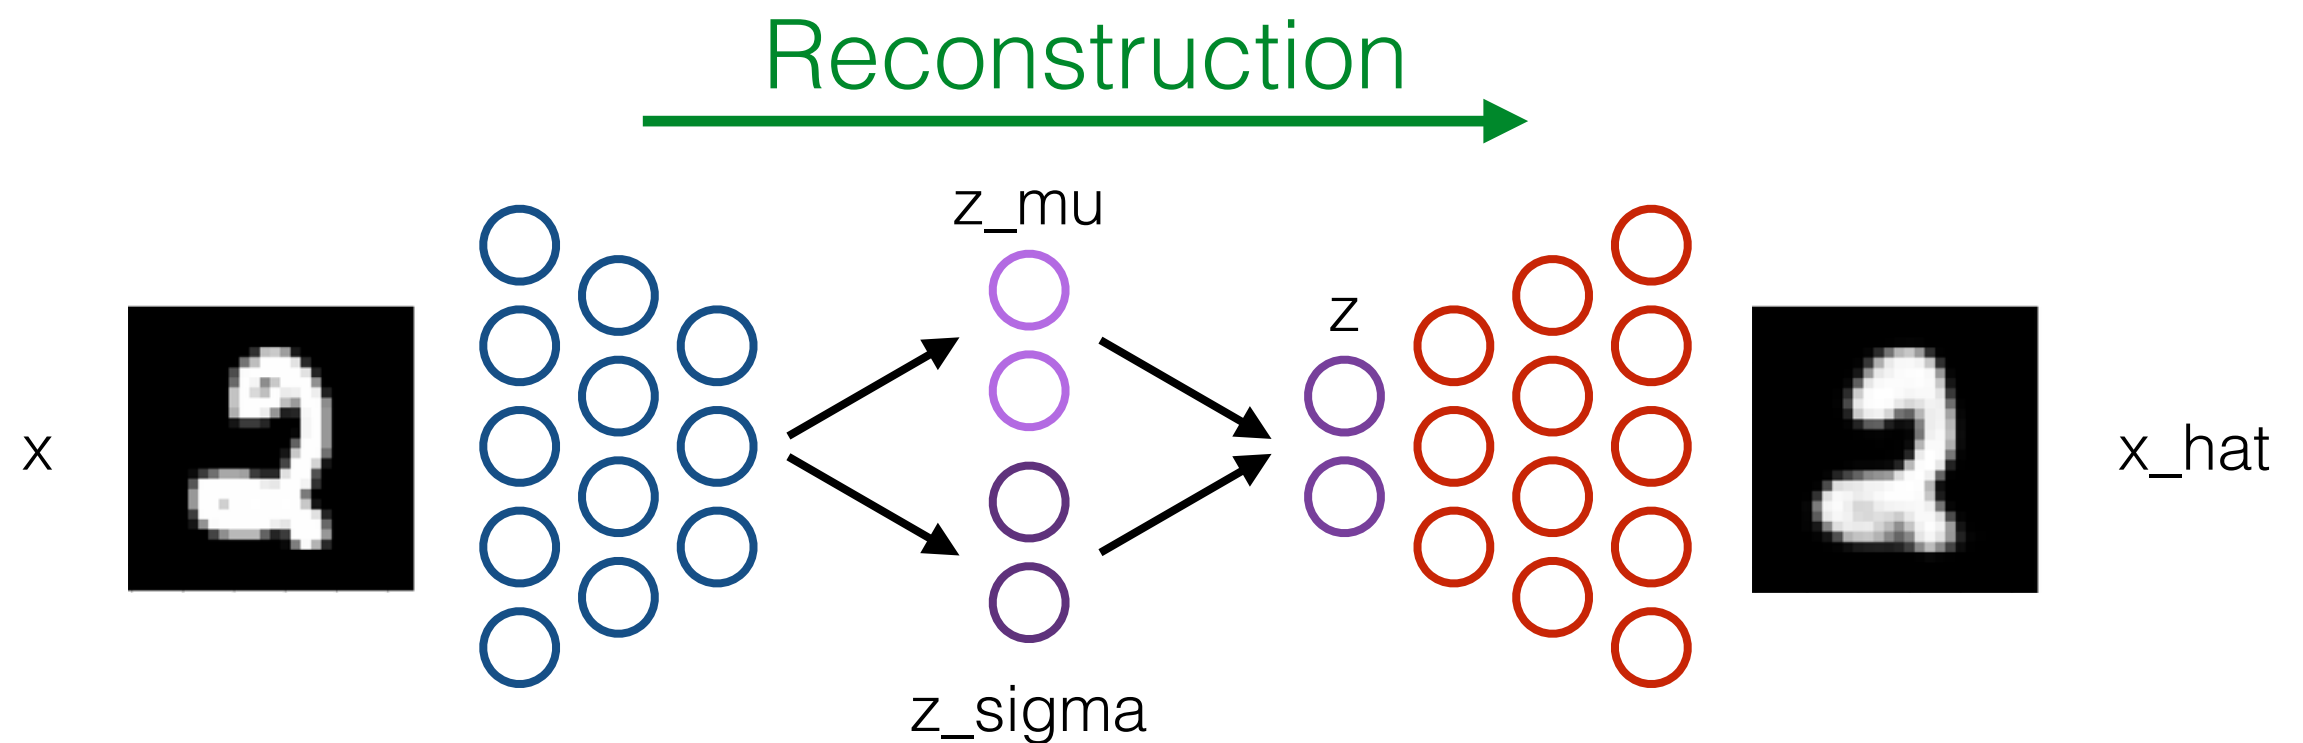
Variational Autoencoder

- Function: training phase
- Train the model given a batch



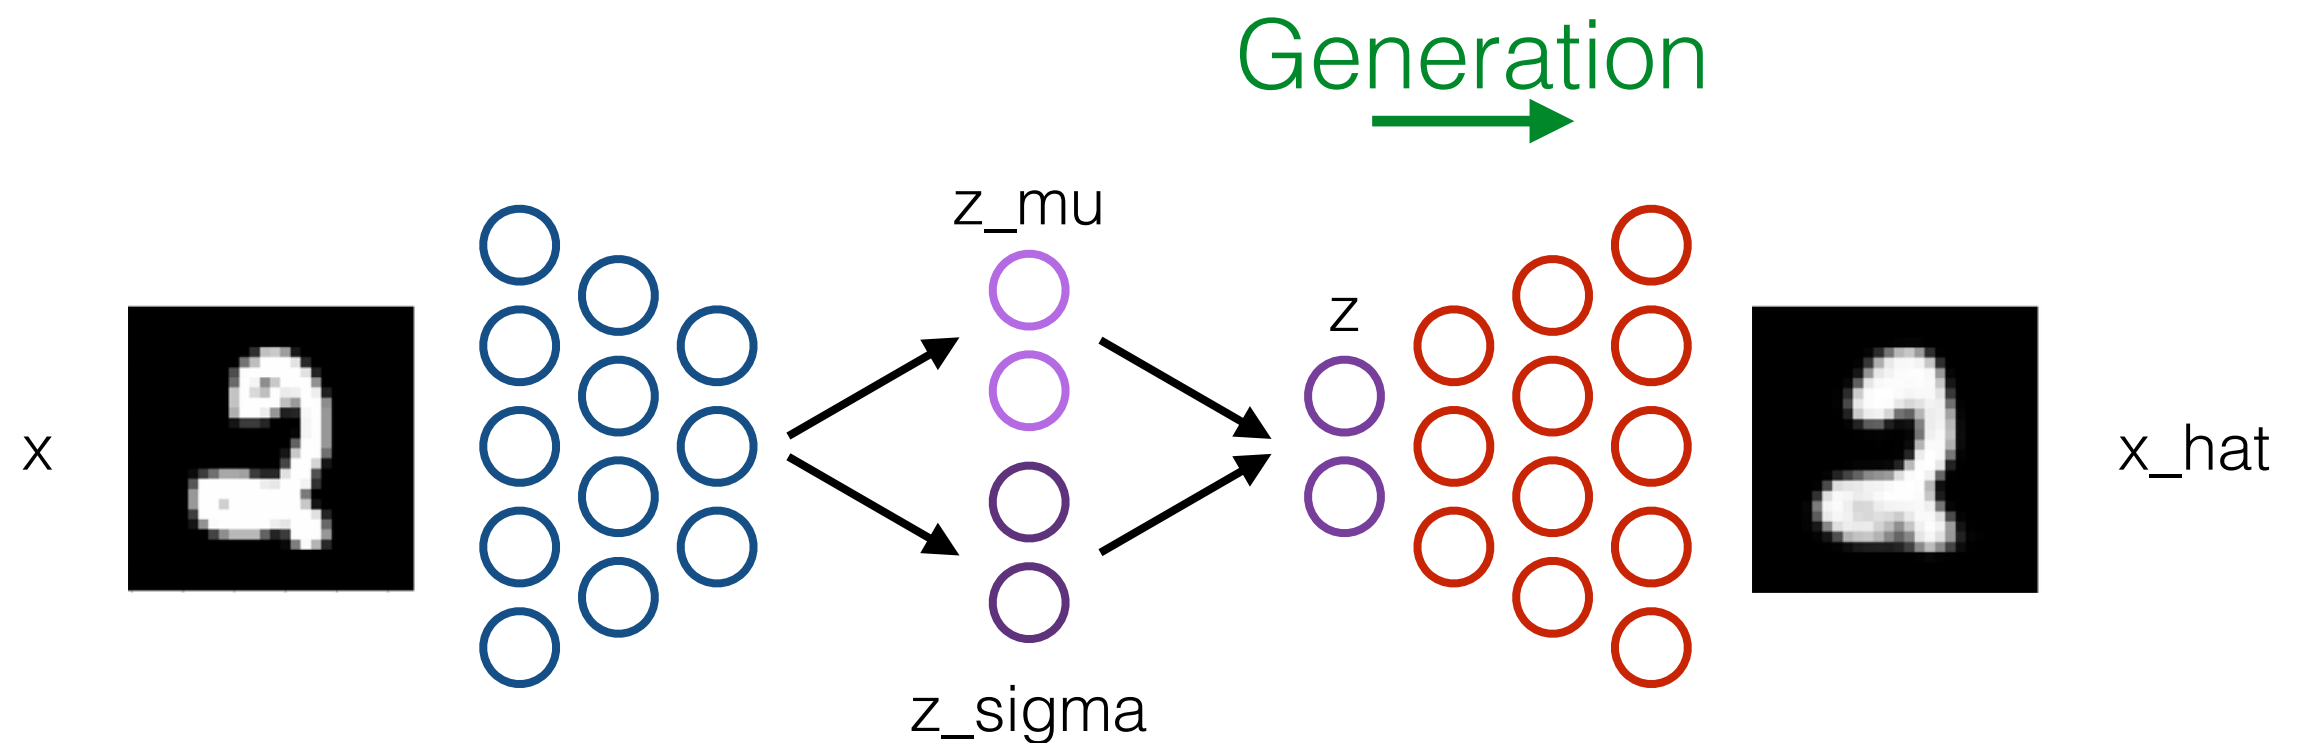
Variational Autoencoder

- Function: testing phase
- Reconstruction



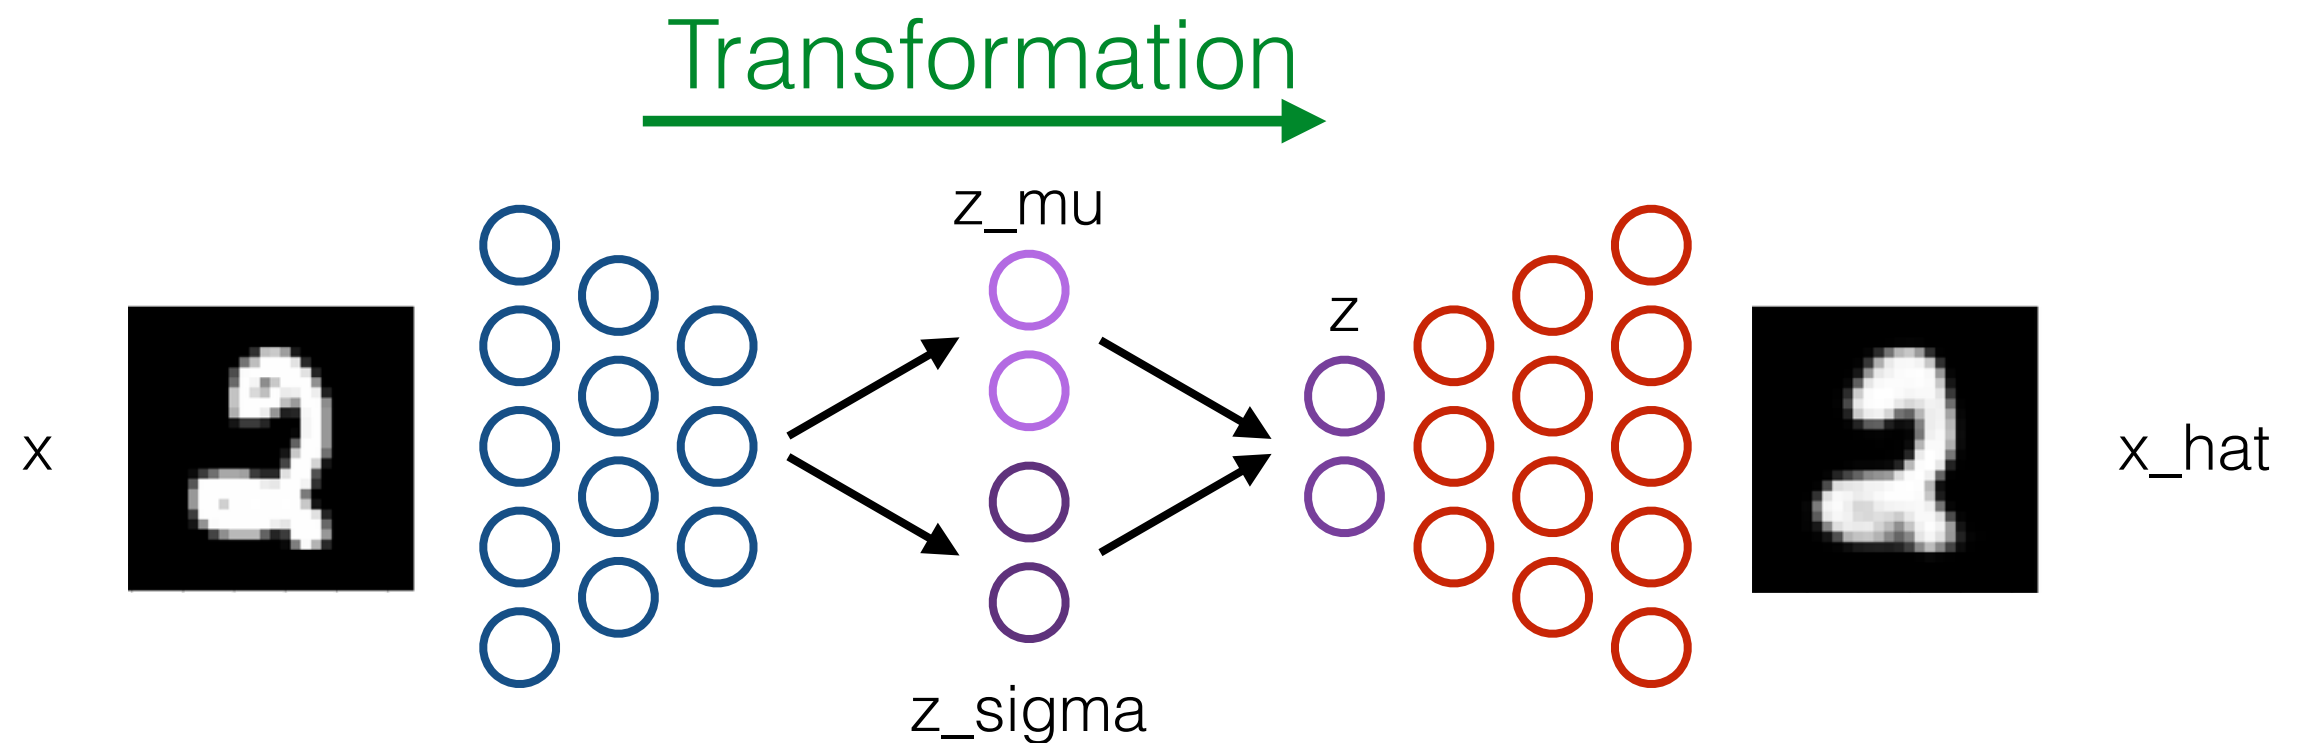
Variational Autoencoder

- Function: testing phase
- Generation



Variational Autoencoder

- Function: testing phase
- Transformation



Coding Session

Variational Autoencoder

Review

Import

```
from tensorflow.examples.tutorials.mnist import input_data
mnist = input_data.read_data_sets('MNIST_data', one_hot=True)
num_sample = mnist.train.num_examples
input_dim = mnist.train.images[0].shape[0]
w = h = int(np.sqrt(input_dim))
```

Data

```
import time
import numpy as np
import tensorflow as tf
from tensorflow.contrib.slim import fully_connected as fc
import matplotlib.pyplot as plt
%matplotlib inline
```

Review

Model

```
class VariationalAutoencoder(object):  
  
    def __init__(self, learning_rate=1e-4, batch_size=64, n_z=16):  
        self.learning_rate = learning_rate  
        self.batch_size = batch_size  
        self.n_z = n_z  
  
        tf.reset_default_graph()  
        self.build()  
  
        self.sess = tf.InteractiveSession()  
        self.sess.run(tf.global_variables_initializer())
```


Review

Build the network: the encoder and decoder

```
# Build the network and the loss functions
def build(self):
    self.x = tf.placeholder(name='x', dtype=tf.float32, shape=[None, input_dim])

    # Encode
    # x -> z_mean, z_sigma -> z
    f1 = fc(self.x, 256, scope='enc_fc1', activation_fn=tf.nn.elu)
    f2 = fc(f1, 128, scope='enc_fc2', activation_fn=tf.nn.elu)
    f3 = fc(f2, 64, scope='enc_fc3', activation_fn=tf.nn.elu)
    self.z_mu = fc(f3, self.n_z, scope='enc_fc4_mu', activation_fn=None)
    self.z_log_sigma_sq = fc(f3, self.n_z, scope='enc_fc4_sigma', activation_fn=None)
    eps = tf.random_normal(
        shape=tf.shape(self.z_log_sigma_sq),
        mean=0, stddev=1, dtype=tf.float32)
    self.z = self.z_mu + tf.sqrt(tf.exp(self.z_log_sigma_sq)) * eps

    # Decode
    # z -> x_hat
    g1 = fc(self.z, 64, scope='dec_fc1', activation_fn=tf.nn.elu)
    g2 = fc(g1, 128, scope='dec_fc2', activation_fn=tf.nn.elu)
    g3 = fc(g2, 256, scope='dec_fc3', activation_fn=tf.nn.elu)
    self.x_hat = fc(g3, input_dim, scope='dec_fc4', activation_fn=tf.sigmoid)
```

Review

Build the network: the loss

```
# Loss
# Reconstruction loss
# Minimize the cross-entropy loss
#  $H(x, x\_hat) = -\sum x \log(x\_hat) + (1-x) \log(1-x\_hat)$ 
epsilon = 1e-10
recon_loss = -tf.reduce_sum(
    self.x * tf.log(epsilon+self.x_hat) + (1-self.x) * tf.log(epsilon+1-self.x_hat),
    axis=1
)
self.recon_loss = tf.reduce_mean(recon_loss)

# Latent loss
# KL divergence: measure the difference between two distributions
# Here we measure the divergence between the latent distribution and  $N(0, 1)$ 
latent_loss = -0.5 * tf.reduce_sum(
    1 + self.z_log_sigma_sq - tf.square(self.z_mu) - tf.exp(self.z_log_sigma_sq), axis=1)
self.latent_loss = tf.reduce_mean(latent_loss)

self.total_loss = self.recon_loss + self.latent_loss
self.train_op = tf.train.AdamOptimizer(
    learning_rate=self.learning_rate).minimize(self.total_loss)

self.losses = {
    'recon_loss': self.recon_loss,
    'latent_loss': self.latent_loss,
    'total_loss': self.total_loss,
}
```

Review

Training and testing functions

```
# Execute the forward and the backward pass
def run_single_step(self, x):
    _, losses = self.sess.run(
        [self.train_op, self.losses],
        feed_dict={self.x: x}
    )
    return losses

# x -> x_hat
def reconstructor(self, x):
    x_hat = self.sess.run(self.x_hat, feed_dict={self.x: x})
    return x_hat

# z -> x
def generator(self, z):
    x_hat = self.sess.run(self.x_hat, feed_dict={self.z: z})
    return x_hat

# x -> z
def transformer(self, x):
    z = self.sess.run(self.z, feed_dict={self.x: x})
    return z
```

Review

Train a model

```
def trainer(model_object, learning_rate=1e-4, batch_size=64, num_epoch=100, n_z=16):
    model = model_object(
        learning_rate=learning_rate, batch_size=batch_size, n_z=n_z)

    for epoch in range(num_epoch):
        start_time = time.time()
        for iter in range(num_sample // batch_size):
            # Get a batch
            batch = mnist.train.next_batch(batch_size)
            # Execute the forward and backward pass and report computed losses
            losses = model.run_single_step(batch[0])
            end_time = time.time()

            if epoch % 5 == 0:
                log_str = '[Epoch {}] '.format(epoch)
                for k, v in losses.items():
                    log_str += '{}: {:.3f} '.format(k, v)
                log_str += '({:.3f} sec/epoch)'.format(end_time - start_time)
                print(log_str)

        print('Done!')
    return model

# Train a model
model = trainer(VariationalAutoencoder)
```


Review

Train a model

```
[Epoch 0] recon_loss: 182.894 latent_loss: 9.692 total_loss: 192.586 (2.595 sec/epoch)
[Epoch 5] recon_loss: 113.784 latent_loss: 14.704 total_loss: 128.488 (2.470 sec/epoch)
[Epoch 10] recon_loss: 108.452 latent_loss: 17.072 total_loss: 125.523 (2.724 sec/epoch)
[Epoch 15] recon_loss: 101.748 latent_loss: 17.579 total_loss: 119.328 (2.027 sec/epoch)
[Epoch 20] recon_loss: 91.644 latent_loss: 18.083 total_loss: 109.727 (2.467 sec/epoch)
[Epoch 25] recon_loss: 94.685 latent_loss: 18.548 total_loss: 113.234 (2.504 sec/epoch)
[Epoch 30] recon_loss: 87.946 latent_loss: 18.468 total_loss: 106.414 (2.454 sec/epoch)
[Epoch 35] recon_loss: 90.971 latent_loss: 18.891 total_loss: 109.862 (2.588 sec/epoch)
[Epoch 40] recon_loss: 84.645 latent_loss: 18.783 total_loss: 103.428 (2.144 sec/epoch)
[Epoch 45] recon_loss: 85.681 latent_loss: 19.136 total_loss: 104.818 (2.589 sec/epoch)
[Epoch 50] recon_loss: 84.453 latent_loss: 19.142 total_loss: 103.595 (2.524 sec/epoch)
[Epoch 55] recon_loss: 86.023 latent_loss: 19.744 total_loss: 105.767 (2.755 sec/epoch)
[Epoch 60] recon_loss: 84.479 latent_loss: 19.800 total_loss: 104.280 (2.405 sec/epoch)
[Epoch 65] recon_loss: 86.094 latent_loss: 19.502 total_loss: 105.596 (2.068 sec/epoch)
[Epoch 70] recon_loss: 84.493 latent_loss: 19.866 total_loss: 104.359 (2.738 sec/epoch)
[Epoch 75] recon_loss: 83.497 latent_loss: 19.736 total_loss: 103.233 (2.659 sec/epoch)
[Epoch 80] recon_loss: 84.940 latent_loss: 20.300 total_loss: 105.240 (2.567 sec/epoch)
[Epoch 85] recon_loss: 83.345 latent_loss: 19.918 total_loss: 103.263 (2.544 sec/epoch)
[Epoch 90] recon_loss: 83.698 latent_loss: 20.042 total_loss: 103.740 (2.691 sec/epoch)
[Epoch 95] recon_loss: 84.776 latent_loss: 19.934 total_loss: 104.710 (2.520 sec/epoch)
Done!
```


Review

Test the model: **reconstruction**

```
def test_reconstruction(model, mnist, h=28, w=28, batch_size=100):  
    # Test the trained model: reconstruction  
    batch = mnist.test.next_batch(batch_size)  
    x_reconstructed = model.reconstructor(batch[0])  
  
    n = np.sqrt(batch_size).astype(np.int32)  
    I_reconstructed = np.empty((h*n, 2*w*n))  
    for i in range(n):  
        for j in range(n):  
            x = np.concatenate(  
                (x_reconstructed[i*n+j, :].reshape(h, w),  
                 batch[0][i*n+j, :].reshape(h, w)),  
                axis=1  
            )  
            I_reconstructed[i*h:(i+1)*h, j*2*w:(j+1)*2*w] = x  
  
    plt.figure(figsize=(10, 20))  
    plt.imshow(I_reconstructed, cmap='gray')
```

Review

Test the model: **reconstruction**



Review

Test the model: **generation**


```
# Test the trained model: generation
# Sample noise vectors from  $N(0, 1)$ 
z = np.random.normal(size=[model.batch_size, model.n_z])
x_generated = model.generator(z)

n = np.sqrt(model.batch_size).astype(np.int32)
I_generated = np.empty((h*n, w*n))
for i in range(n):
    for j in range(n):
        I_generated[i*h:(i+1)*h, j*w:(j+1)*w] = x_generated[i*n+j, :].reshape(h, w)

plt.figure(figsize=(8, 8))
plt.imshow(I_generated, cmap='gray')
```


Review

Test the model: **generation**



Review

Train a 2d model

```
# Train a model with 2d latent space  
model_2d = trainer(VariationalAutoencoder, n_z=2)
```

```
[Epoch 0] recon_loss: 199.729 latent_loss: 4.342 total_loss: 204.071 (2.159 sec/epoch)  
[Epoch 5] recon_loss: 167.280 latent_loss: 4.728 total_loss: 172.008 (2.601 sec/epoch)  
[Epoch 10] recon_loss: 145.253 latent_loss: 5.516 total_loss: 150.769 (2.589 sec/epoch)  
[Epoch 15] recon_loss: 144.056 latent_loss: 5.380 total_loss: 149.436 (2.580 sec/epoch)  
[Epoch 20] recon_loss: 148.198 latent_loss: 5.860 total_loss: 154.058 (2.712 sec/epoch)  
[Epoch 25] recon_loss: 152.884 latent_loss: 5.783 total_loss: 158.667 (2.075 sec/epoch)  
[Epoch 30] recon_loss: 147.188 latent_loss: 5.782 total_loss: 152.970 (2.813 sec/epoch)  
[Epoch 35] recon_loss: 144.544 latent_loss: 5.971 total_loss: 150.515 (2.479 sec/epoch)  
[Epoch 40] recon_loss: 152.908 latent_loss: 5.965 total_loss: 158.873 (2.542 sec/epoch)  
[Epoch 45] recon_loss: 132.893 latent_loss: 6.022 total_loss: 138.916 (2.746 sec/epoch)  
[Epoch 50] recon_loss: 136.282 latent_loss: 6.303 total_loss: 142.585 (2.093 sec/epoch)  
[Epoch 55] recon_loss: 148.406 latent_loss: 6.235 total_loss: 154.640 (2.798 sec/epoch)  
[Epoch 60] recon_loss: 131.068 latent_loss: 6.062 total_loss: 137.130 (2.410 sec/epoch)  
[Epoch 65] recon_loss: 135.531 latent_loss: 6.250 total_loss: 141.781 (2.393 sec/epoch)  
[Epoch 70] recon_loss: 129.911 latent_loss: 6.276 total_loss: 136.187 (2.722 sec/epoch)  
[Epoch 75] recon_loss: 141.734 latent_loss: 6.252 total_loss: 147.986 (2.587 sec/epoch)  
[Epoch 80] recon_loss: 149.359 latent_loss: 6.349 total_loss: 155.708 (2.023 sec/epoch)  
[Epoch 85] recon_loss: 138.324 latent_loss: 6.197 total_loss: 144.521 (2.484 sec/epoch)  
[Epoch 90] recon_loss: 130.314 latent_loss: 6.396 total_loss: 136.711 (2.500 sec/epoch)  
[Epoch 95] recon_loss: 133.127 latent_loss: 6.579 total_loss: 139.706 (2.602 sec/epoch)  
Done!
```


Review

Test the 2d model: **transformation**

```
def test_transformation(model_2d, mnist, batch_size=3000):  
    # Test the trained model: transformation  
    assert model.n_z == 2  
    batch = mnist.test.next_batch(batch_size)  
    z = model_2d.transformer(batch[0])  
    plt.figure(figsize=(10, 8))  
    plt.scatter(z[:, 0], z[:, 1], c=np.argmax(batch[1], 1), s=20)  
    plt.colorbar()  
    plt.grid()
```

Review

Test the 2d model: **transformation**



Review

Test the 2d model: **latent space**

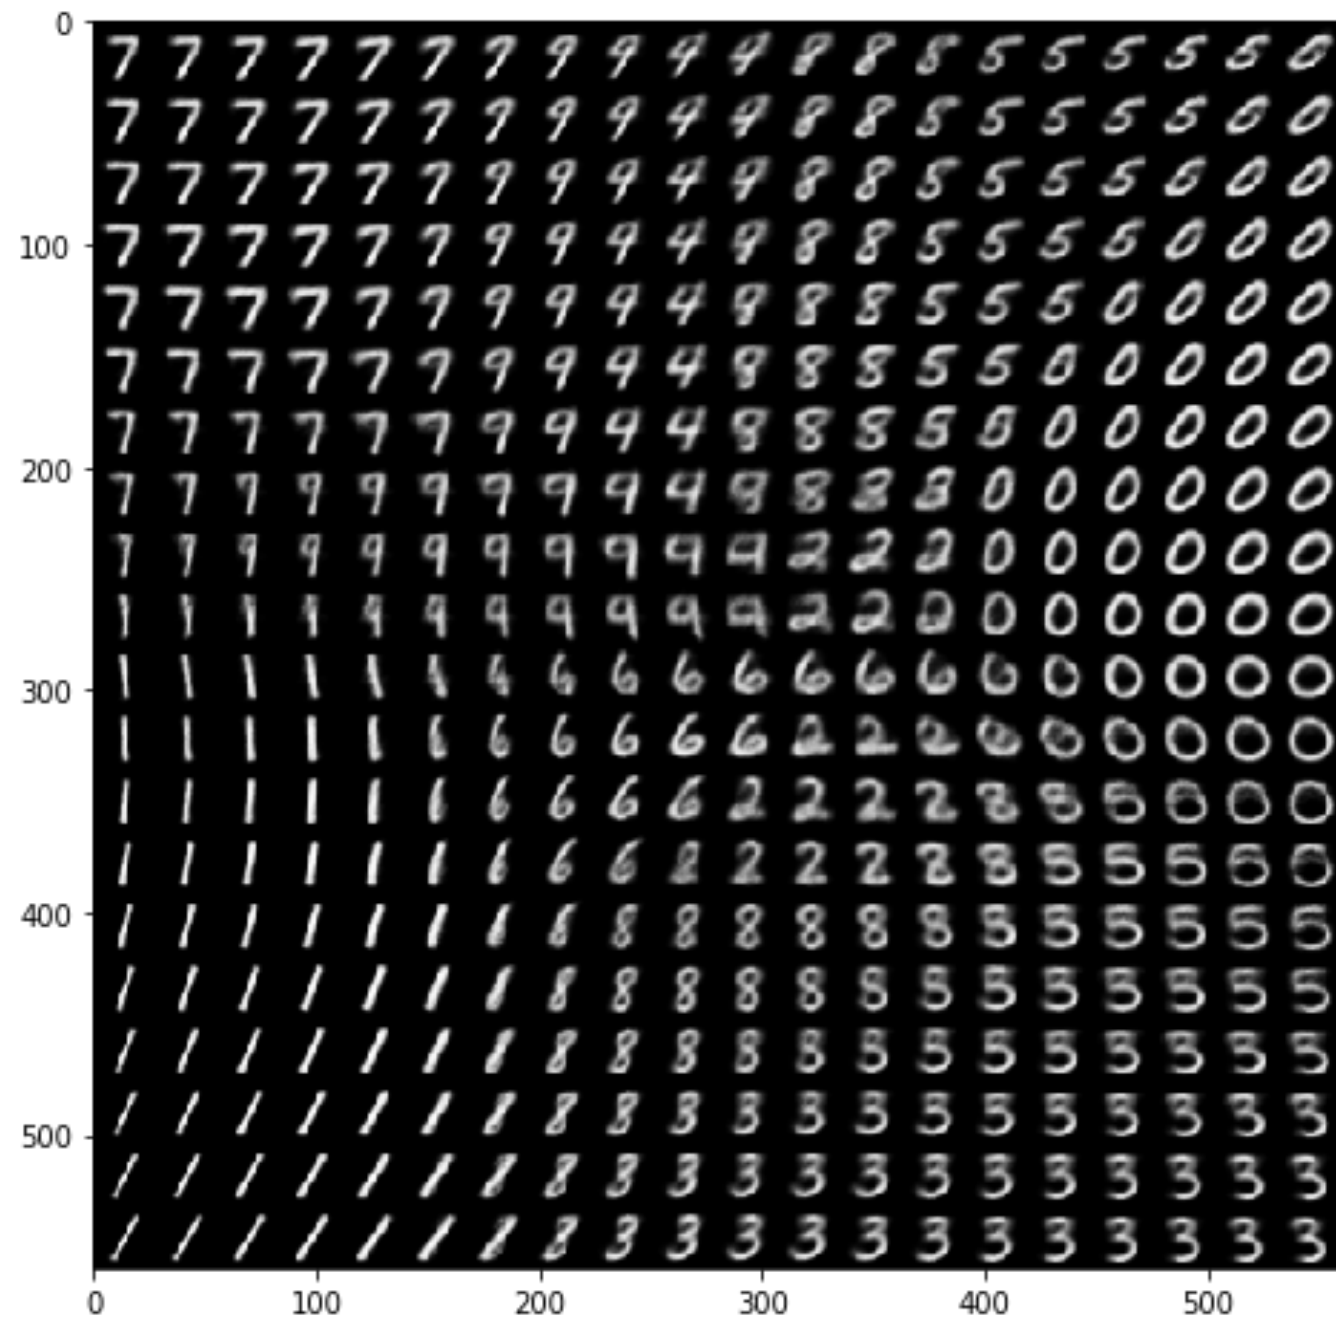
```
# Test the trained model: transformation
n = 20
x = np.linspace(-2, 2, n)
y = np.linspace(-2, 2, n)

I_latent = np.empty((h*n, w*n))
for i, yi in enumerate(x):
    for j, xi in enumerate(y):
        z = np.array([[xi, yi]]*model_2d.batch_size)
        x_hat = model_2d.generator(z)
        I_latent[(n-i-1)*h:(n-i)*h, j*w:(j+1)*w] = x_hat[0].reshape(h, w)

plt.figure(figsize=(8, 8))
plt.imshow(I_latent, cmap="gray")
```



Review

Test the 2d model: **latent space**




Review

Test the 2d model: **latent space**



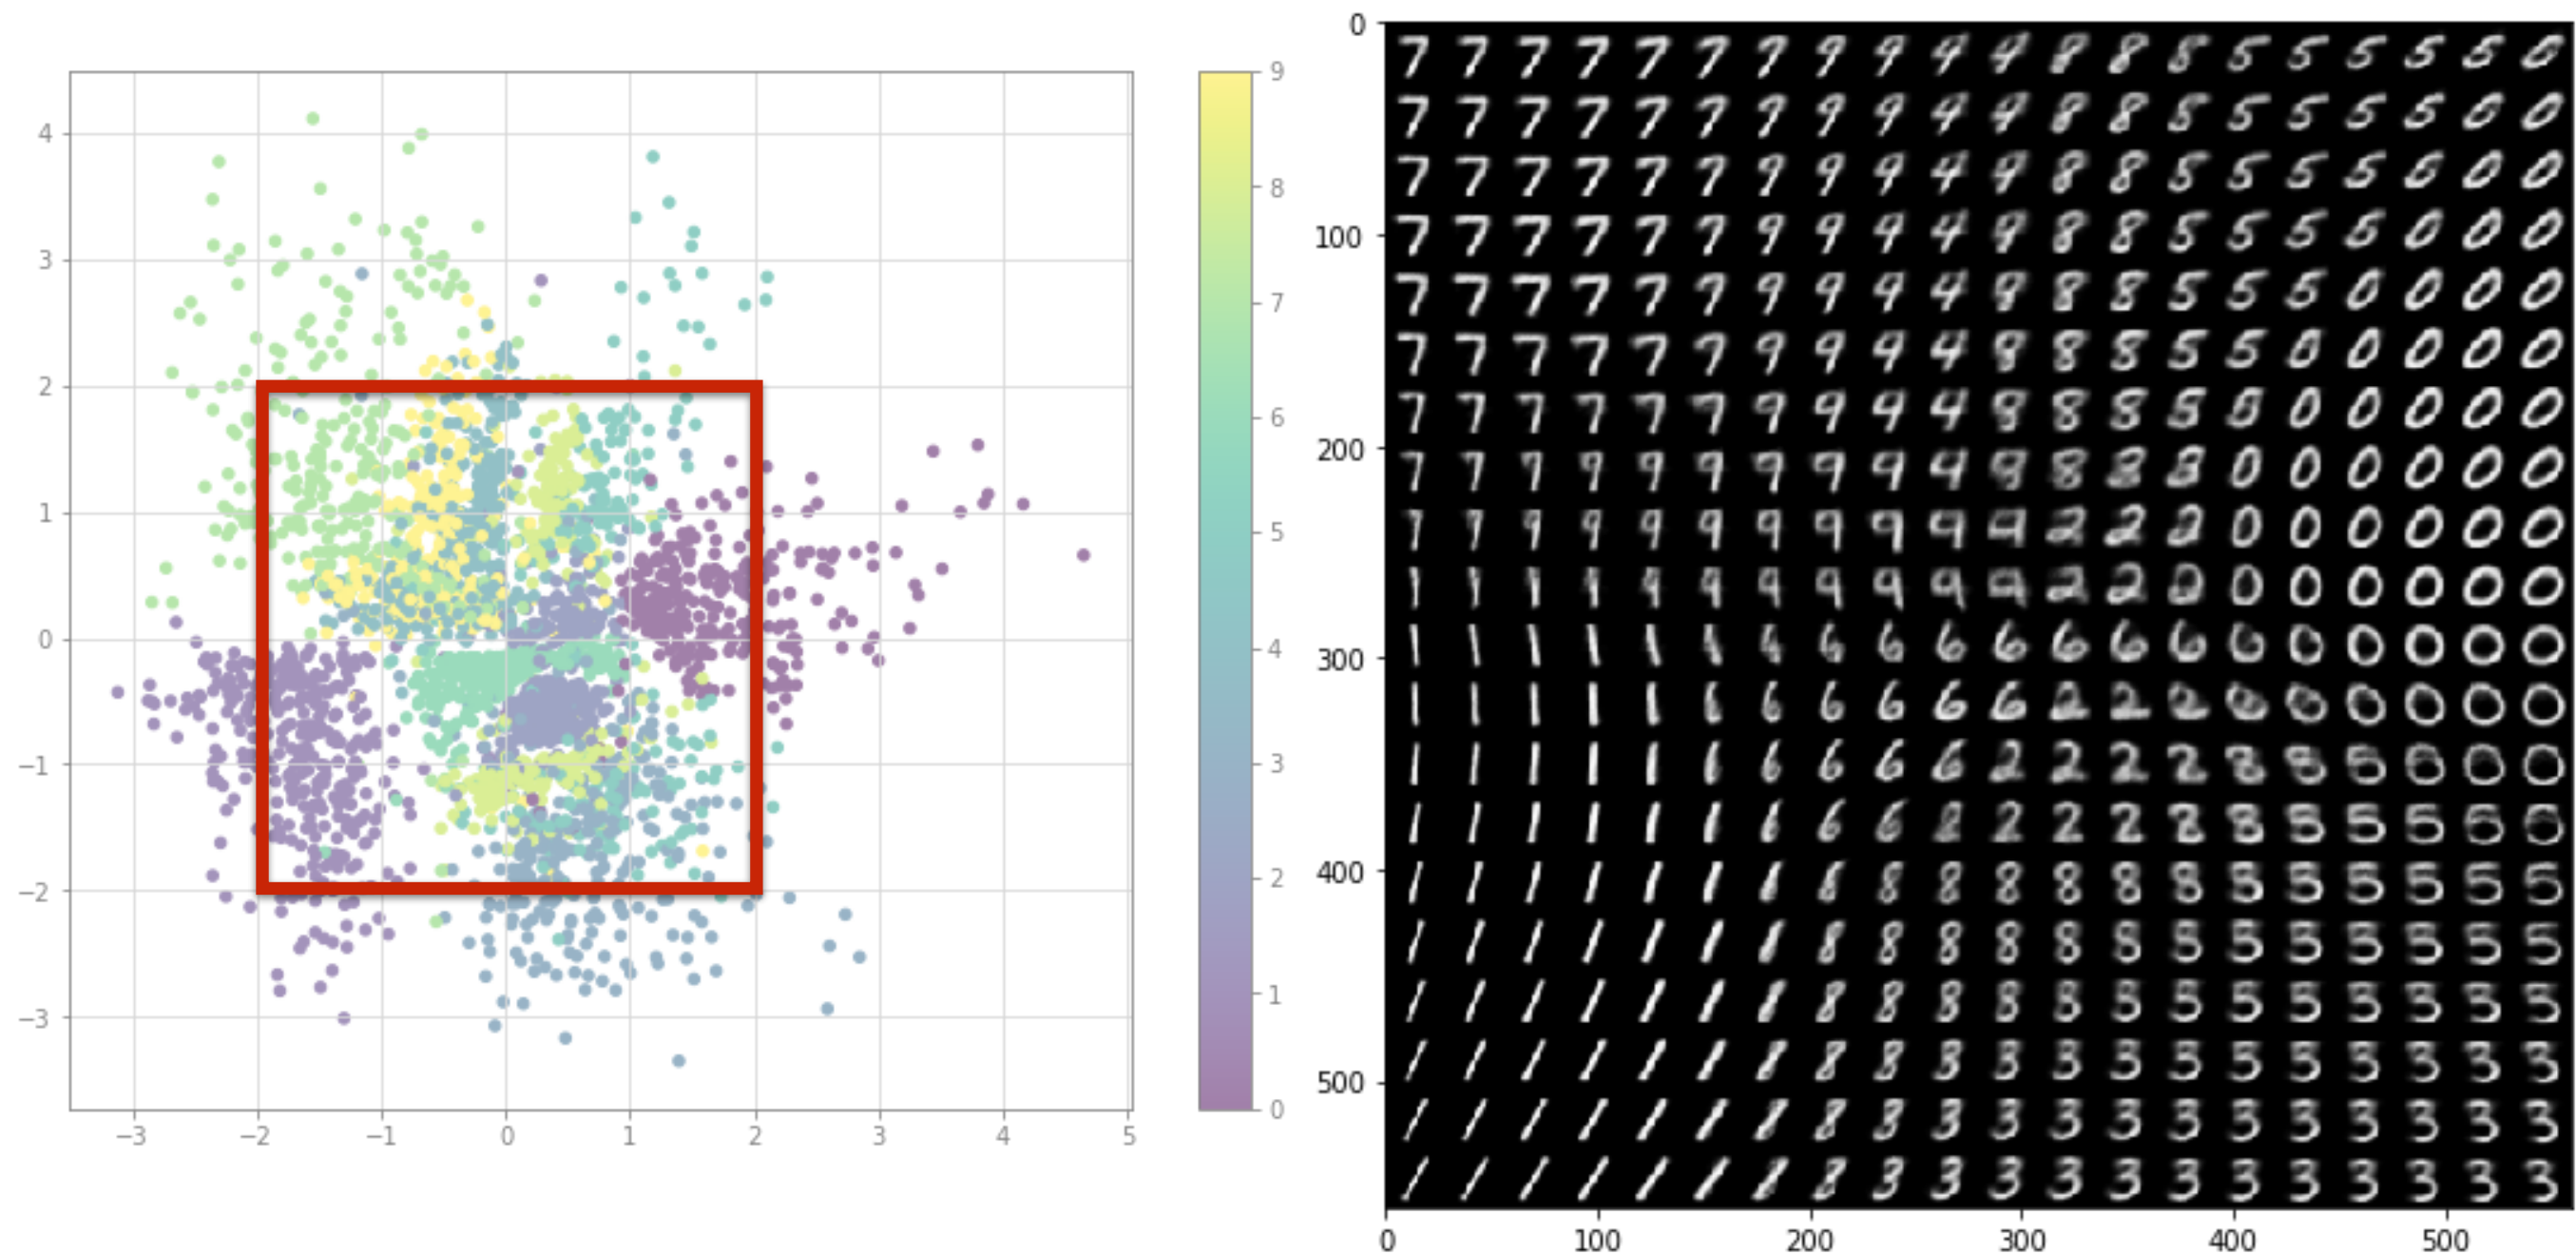
Review

Test the 2d model: **latent space**



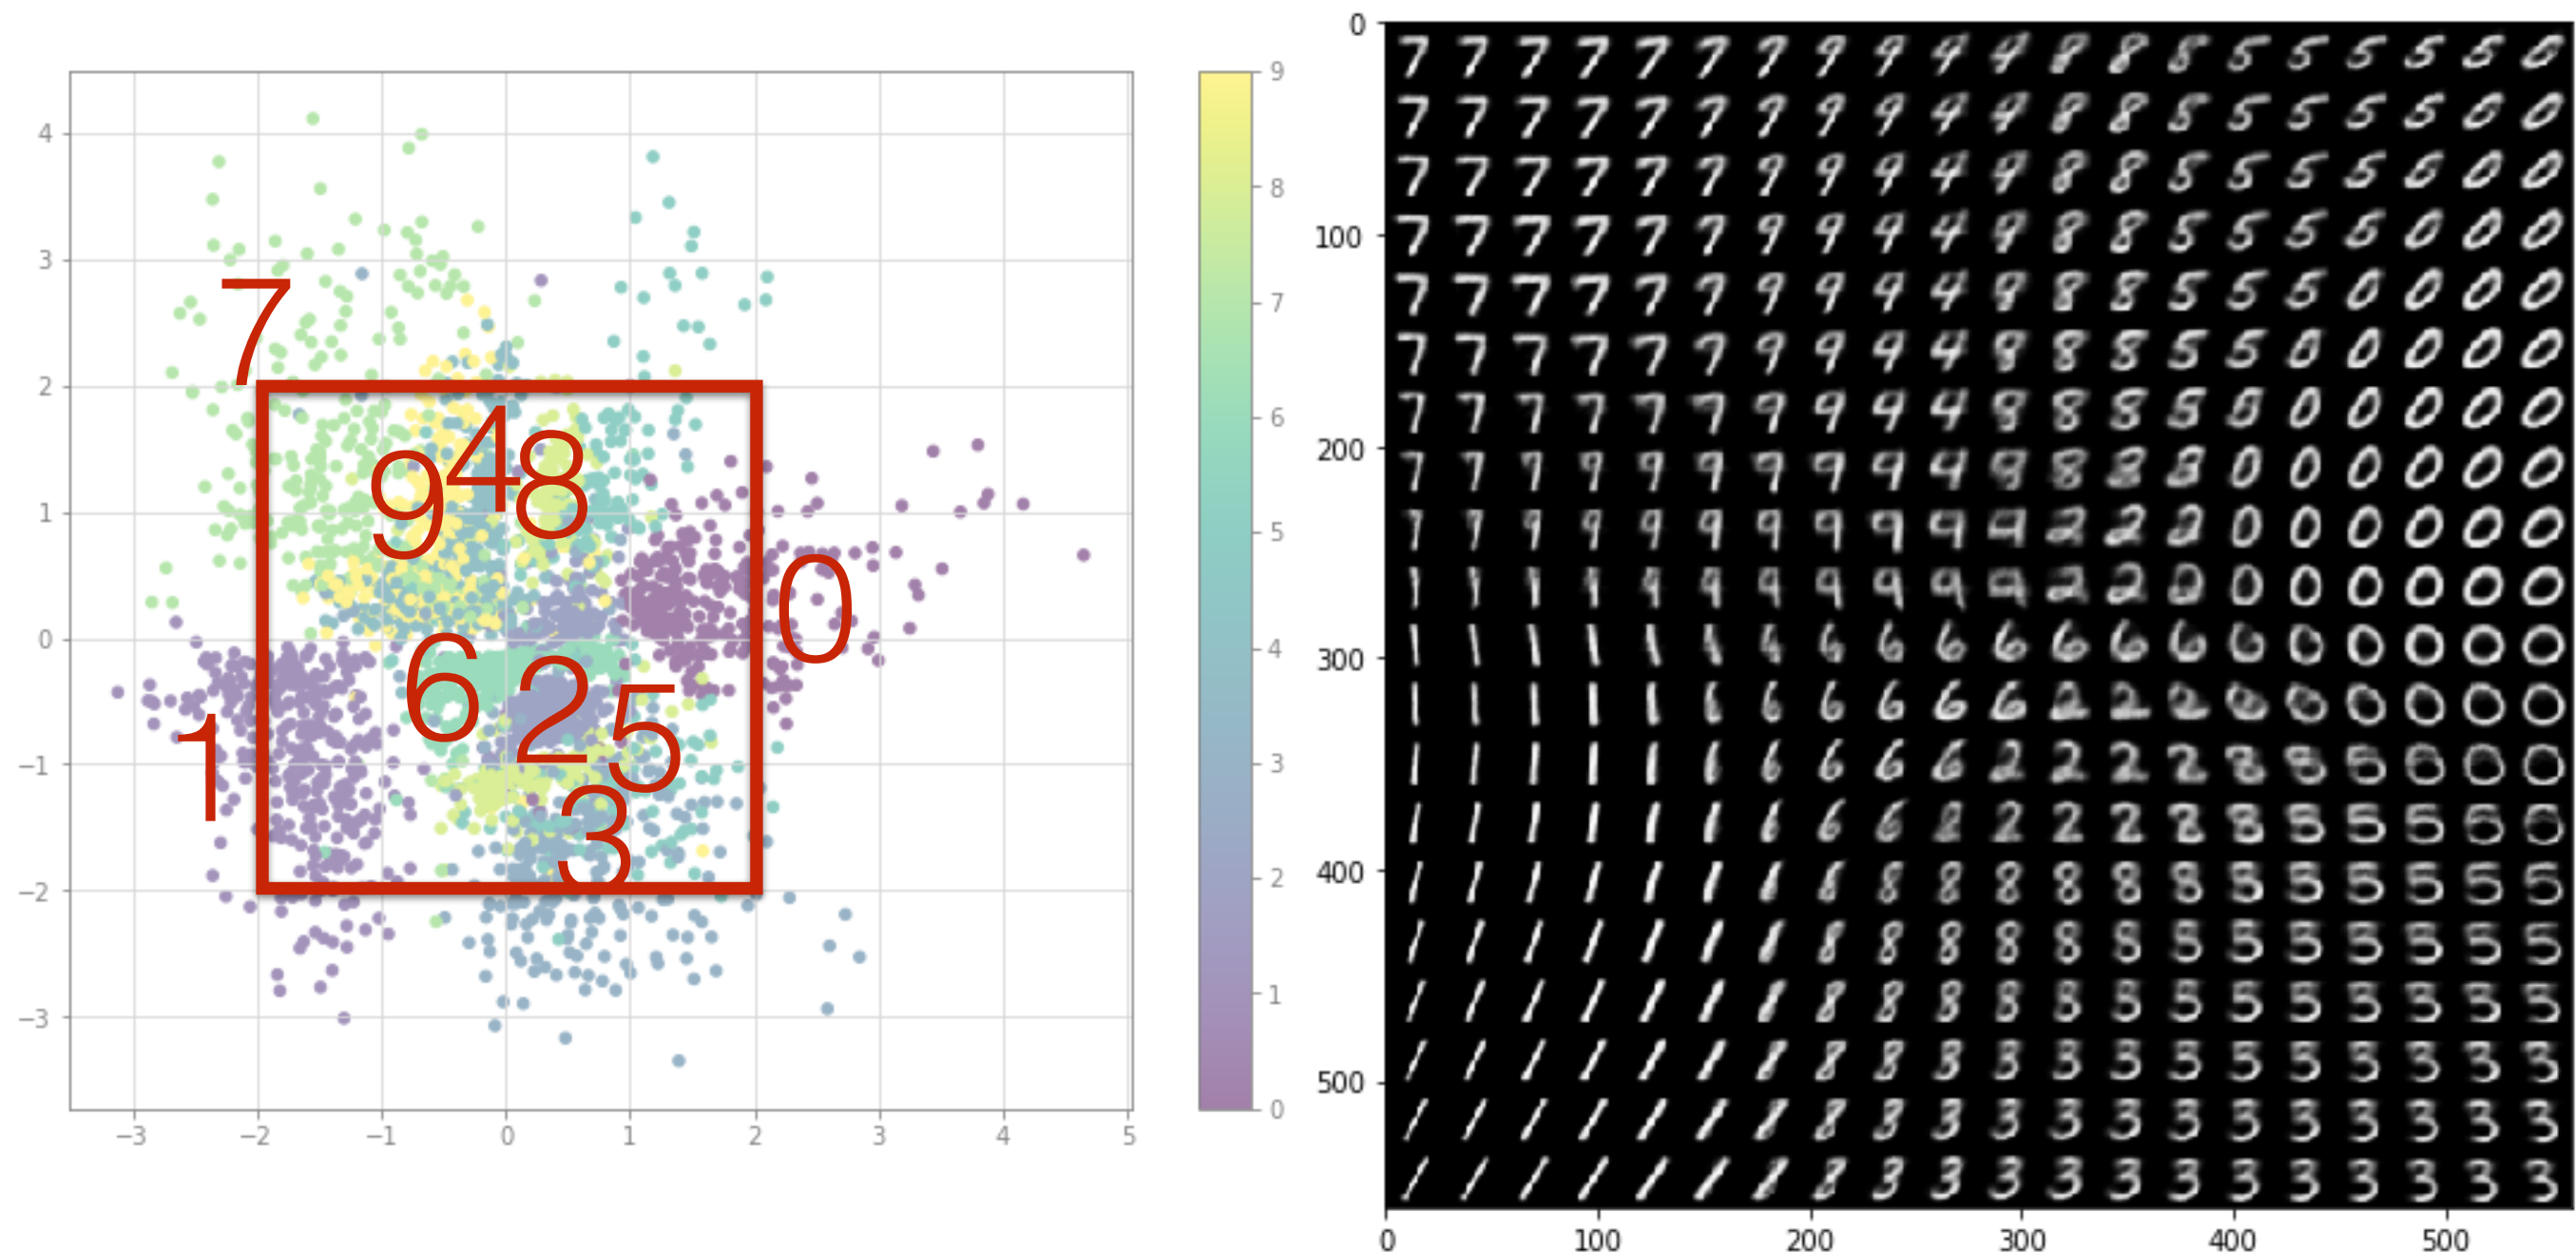
Review

Test the 2d model: **latent space**



Review

Test the 2d model: **latent space**



Summary

TensorFlow

- Build the graph
- Run the graph

Variational Autoencoder

- Generative model
- [A blog post about VAEs](#)
- Demo code (<https://github.com/shaohua0116/VAE-Tensorflow>)

Questions?

Today's agenda

- **Part 1: Deep Learning Framework**
 - TensorFlow
 - PyTorch
- **Part 2: Cloud Service**
 - Google Cloud
 - Amazon Web Services